

# 11-2-linear-regression-1

April 27, 2024

## 1 Seatwork 11.1 Exploratory Data Analysis for Machine Learning

```
[1]: pip install ucimlrepo
```

Collecting ucimlrepo

Downloading ucimlrepo-0.0.6-py3-none-any.whl (8.0 kB)

Installing collected packages: ucimlrepo

Successfully installed ucimlrepo-0.0.6

```
[2]: from ucimlrepo import fetch_ucirepo
# fetch dataset
automobile = fetch_ucirepo(id=10)
# data (as pandas dataframes)
X = automobile.data.features
y = automobile.data.targets
```

```
[3]: X
```

```
[3]:
```

	price	highway-mpg	city-mpg	peak-rpm	horsepower	compression-ratio	\
0	13495.0	27	21	5000.0	111.0	9.0	
1	16500.0	27	21	5000.0	111.0	9.0	
2	16500.0	26	19	5000.0	154.0	9.0	
3	13950.0	30	24	5500.0	102.0	10.0	
4	17450.0	22	18	5500.0	115.0	8.0	
..	...	...	...	...	...	...	
200	16845.0	28	23	5400.0	114.0	9.5	
201	19045.0	25	19	5300.0	160.0	8.7	
202	21485.0	23	18	5500.0	134.0	8.8	
203	22470.0	27	26	4800.0	106.0	23.0	
204	22625.0	25	19	5400.0	114.0	9.5	

	stroke	bore	fuel-system	engine-size	...	length	wheel-base	\
0	2.68	3.47	mpfi	130	...	168.8	88.6	
1	2.68	3.47	mpfi	130	...	168.8	88.6	
2	3.47	2.68	mpfi	152	...	171.2	94.5	
3	3.40	3.19	mpfi	109	...	176.6	99.8	
4	3.40	3.19	mpfi	136	...	176.6	99.4	

```

..      ...      ...      ...      ...      ...      ...
200      3.15  3.78      mpfi      141  ...  188.8      109.1
201      3.15  3.78      mpfi      141  ...  188.8      109.1
202      2.87  3.58      mpfi      173  ...  188.8      109.1
203      3.40  3.01      idi      145  ...  188.8      109.1
204      3.15  3.78      mpfi      141  ...  188.8      109.1

      engine-location  drive-wheels  body-style  num-of-doors  aspiration \
0              front      rwd  convertible      2.0      std
1              front      rwd  convertible      2.0      std
2              front      rwd   hatchback      2.0      std
3              front      fwd    sedan      4.0      std
4              front      4wd    sedan      4.0      std
..              ...      ...      ...      ...      ...
200              front      rwd    sedan      4.0      std
201              front      rwd    sedan      4.0     turbo
202              front      rwd    sedan      4.0      std
203              front      rwd    sedan      4.0     turbo
204              front      rwd    sedan      4.0     turbo

      fuel-type      make  normalized-losses
0          gas  alfa-romero      NaN
1          gas  alfa-romero      NaN
2          gas  alfa-romero      NaN
3          gas      audi     164.0
4          gas      audi     164.0
..              ...      ...      ...
200          gas      volvo     95.0
201          gas      volvo     95.0
202          gas      volvo     95.0
203      diesel      volvo     95.0
204          gas      volvo     95.0

```

[205 rows x 25 columns]

[4]: y

```

[4]:      symboling
0          3
1          3
2          1
3          2
4          2
..      ...
200      -1
201      -1
202      -1

```

```
203         -1
204         -1
```

```
[205 rows x 1 columns]
```

```
[6]: !pip install hvplot
```

```
Collecting hvplot
```

```
  Downloading hvplot-0.9.2-py2.py3-none-any.whl (1.8 MB)
```

```
      1.8/1.8 MB
```

```
8.2 MB/s eta 0:00:00
```

```
Requirement already satisfied: bokeh>=1.0.0 in
```

```
/usr/local/lib/python3.10/dist-packages (from hvplot) (3.3.4)
```

```
Requirement already satisfied: colorcet>=2 in /usr/local/lib/python3.10/dist-  
packages (from hvplot) (3.1.0)
```

```
Requirement already satisfied: holoviews>=1.11.0 in
```

```
/usr/local/lib/python3.10/dist-packages (from hvplot) (1.17.1)
```

```
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages  
(from hvplot) (2.0.3)
```

```
Requirement already satisfied: numpy>=1.15 in /usr/local/lib/python3.10/dist-  
packages (from hvplot) (1.25.2)
```

```
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-  
packages (from hvplot) (24.0)
```

```
Requirement already satisfied: panel>=0.11.0 in /usr/local/lib/python3.10/dist-  
packages (from hvplot) (1.3.8)
```

```
Requirement already satisfied: param<3.0,>=1.12.0 in
```

```
/usr/local/lib/python3.10/dist-packages (from hvplot) (2.1.0)
```

```
Requirement already satisfied: Jinja2>=2.9 in /usr/local/lib/python3.10/dist-  
packages (from bokeh>=1.0.0->hvplot) (3.1.3)
```

```
Requirement already satisfied: contourpy>=1 in /usr/local/lib/python3.10/dist-  
packages (from bokeh>=1.0.0->hvplot) (1.2.1)
```

```
Requirement already satisfied: pillow>=7.1.0 in /usr/local/lib/python3.10/dist-  
packages (from bokeh>=1.0.0->hvplot) (9.4.0)
```

```
Requirement already satisfied: PyYAML>=3.10 in /usr/local/lib/python3.10/dist-  
packages (from bokeh>=1.0.0->hvplot) (6.0.1)
```

```
Requirement already satisfied: tornado>=5.1 in /usr/local/lib/python3.10/dist-  
packages (from bokeh>=1.0.0->hvplot) (6.3.3)
```

```
Requirement already satisfied: xyzservices>=2021.09.1 in
```

```
/usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (2024.4.0)
```

```
Requirement already satisfied: pyviz-comms>=0.7.4 in
```

```
/usr/local/lib/python3.10/dist-packages (from holoviews>=1.11.0->hvplot) (3.0.2)
```

```
Requirement already satisfied: python-dateutil>=2.8.2 in
```

```
/usr/local/lib/python3.10/dist-packages (from pandas->hvplot) (2.8.2)
```

```
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-  
packages (from pandas->hvplot) (2023.4)
```

```
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-  
packages (from pandas->hvplot) (2024.1)
```

Requirement already satisfied: markdown in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (3.6)

Requirement already satisfied: markdown-it-py in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (3.0.0)

Requirement already satisfied: linkify-it-py in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (2.0.3)

Requirement already satisfied: mdit-py-plugins in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (0.4.0)

Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (2.31.0)

Requirement already satisfied: tqdm>=4.48.0 in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (4.66.2)

Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (6.1.0)

Requirement already satisfied: typing-extensions in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (4.11.0)

Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from Jinja2>=2.9->bokeh>=1.0.0->hvplot) (2.1.5)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas->hvplot) (1.16.0)

Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from bleach->panel>=0.11.0->hvplot) (0.5.1)

Requirement already satisfied: uc-micro-py in /usr/local/lib/python3.10/dist-packages (from linkify-it-py->panel>=0.11.0->hvplot) (1.0.3)

Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-it-py->panel>=0.11.0->hvplot) (0.1.2)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (3.3.2)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (3.7)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (2.0.7)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (2024.2.2)

Installing collected packages: hvplot

Successfully installed hvplot-0.9.2

```
[7]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import hvplot.pandas
from sklearn.model_selection import train_test_split
```

```

from sklearn import metrics
from sklearn.linear_model import LinearRegression
%matplotlib inline

```

```

[8]: df = pd.concat([X, y], axis = 1)
df

```

```

[8]:      price  highway-mpg  city-mpg  peak-rpm  horsepower  compression-ratio  \
0    13495.0           27         21    5000.0         111.0             9.0
1    16500.0           27         21    5000.0         111.0             9.0
2    16500.0           26         19    5000.0         154.0             9.0
3    13950.0           30         24    5500.0         102.0            10.0
4    17450.0           22         18    5500.0         115.0             8.0
..      ...           ...         ...      ...      ...           ...
200   16845.0           28         23    5400.0         114.0             9.5
201   19045.0           25         19    5300.0         160.0             8.7
202   21485.0           23         18    5500.0         134.0             8.8
203   22470.0           27         26    4800.0         106.0            23.0
204   22625.0           25         19    5400.0         114.0             9.5

```

```

      stroke  bore fuel-system  engine-size  ...  wheel-base  engine-location  \
0      2.68  3.47      mpfi          130  ...      88.6         front
1      2.68  3.47      mpfi          130  ...      88.6         front
2      3.47  2.68      mpfi          152  ...      94.5         front
3      3.40  3.19      mpfi          109  ...      99.8         front
4      3.40  3.19      mpfi          136  ...      99.4         front
..      ...  ...           ...      ...  ...      ...           ...
200     3.15  3.78      mpfi          141  ...      109.1        front
201     3.15  3.78      mpfi          141  ...      109.1        front
202     2.87  3.58      mpfi          173  ...      109.1        front
203     3.40  3.01      idi           145  ...      109.1        front
204     3.15  3.78      mpfi          141  ...      109.1        front

```

```

      drive-wheels  body-style  num-of-doors  aspiration  fuel-type  \
0              rwd  convertible          2.0         std        gas
1              rwd  convertible          2.0         std        gas
2              rwd   hatchback          2.0         std        gas
3              fwd     sedan          4.0         std        gas
4              4wd     sedan          4.0         std        gas
..              ...           ...           ...      ...      ...
200             rwd     sedan          4.0         std        gas
201             rwd     sedan          4.0        turbo        gas
202             rwd     sedan          4.0         std        gas
203             rwd     sedan          4.0        turbo     diesel
204             rwd     sedan          4.0        turbo        gas

```

make normalized-losses symboling

0	alfa-romero	NaN	3
1	alfa-romero	NaN	3
2	alfa-romero	NaN	1
3	audi	164.0	2
4	audi	164.0	2
..	...	...	...
200	volvo	95.0	-1
201	volvo	95.0	-1
202	volvo	95.0	-1
203	volvo	95.0	-1
204	volvo	95.0	-1

[205 rows x 26 columns]

## 2 Cleaning

```
[9]: missing_values = df.isnull().sum()
print(missing_values)
```

```
price          4
highway-mpg    0
city-mpg       0
peak-rpm       2
horsepower     2
compression-ratio 0
stroke         4
bore           4
fuel-system    0
engine-size    0
num-of-cylinders 0
engine-type    0
curb-weight    0
height         0
width          0
length         0
wheel-base    0
engine-location 0
drive-wheels   0
body-style     0
num-of-doors   2
aspiration     0
fuel-type      0
make           0
normalized-losses 41
symboling      0
dtype: int64
```

```
[14]: null_columns = []
      for x in df.columns:
          if df[x].isnull().any():
              null_columns.append(x)
```

```
[15]: null_columns
```

```
[15]: ['price',
      'peak-rpm',
      'horsepower',
      'stroke',
      'bore',
      'num-of-doors',
      'normalized-losses']
```

```
[19]: for x in null_columns:
      df[x] = df[x].fillna(df[x].mean())
```

```
[20]: df
```

```
[20]:      price  highway-mpg  city-mpg  peak-rpm  horsepower  compression-ratio \
0    13495.0           27        21    5000.0         111.0             9.0
1    16500.0           27        21    5000.0         111.0             9.0
2    16500.0           26        19    5000.0         154.0             9.0
3    13950.0           30        24    5500.0         102.0            10.0
4    17450.0           22        18    5500.0         115.0             8.0
..      ...           ...        ...      ...      ...           ...
200   16845.0           28        23    5400.0         114.0             9.5
201   19045.0           25        19    5300.0         160.0             8.7
202   21485.0           23        18    5500.0         134.0             8.8
203   22470.0           27        26    4800.0         106.0            23.0
204   22625.0           25        19    5400.0         114.0             9.5
```

```
      stroke  bore  fuel-system  engine-size  ...  wheel-base  engine-location \
0      2.68  3.47         mpfi          130  ...      88.6         front
1      2.68  3.47         mpfi          130  ...      88.6         front
2      3.47  2.68         mpfi          152  ...      94.5         front
3      3.40  3.19         mpfi          109  ...      99.8         front
4      3.40  3.19         mpfi          136  ...      99.4         front
..      ...  ...           ...      ...      ...           ...
200     3.15  3.78         mpfi          141  ...     109.1         front
201     3.15  3.78         mpfi          141  ...     109.1         front
202     2.87  3.58         mpfi          173  ...     109.1         front
203     3.40  3.01         idi           145  ...     109.1         front
204     3.15  3.78         mpfi          141  ...     109.1         front
```

```
      drive-wheels  body-style  num-of-doors  aspiration  fuel-type \
```

0	rwd	convertible	2.0	std	gas
1	rwd	convertible	2.0	std	gas
2	rwd	hatchback	2.0	std	gas
3	fwd	sedan	4.0	std	gas
4	4wd	sedan	4.0	std	gas
..	...	...	...	...	...
200	rwd	sedan	4.0	std	gas
201	rwd	sedan	4.0	turbo	gas
202	rwd	sedan	4.0	std	gas
203	rwd	sedan	4.0	turbo	diesel
204	rwd	sedan	4.0	turbo	gas

```

make normalized-losses symboling
0   alfa-romero      122.0      3
1   alfa-romero      122.0      3
2   alfa-romero      122.0      1
3       audi         164.0      2
4       audi         164.0      2
..   ...            ...        ...
200  volvo           95.0      -1
201  volvo           95.0      -1
202  volvo           95.0      -1
203  volvo           95.0      -1
204  volvo           95.0      -1

```

[205 rows x 26 columns]

```
[22]: df.isnull().sum()
```

```

[22]: price           0
highway-mpg          0
city-mpg             0
peak-rpm             0
horsepower           0
compression-ratio    0
stroke              0
bore                0
fuel-system          0
engine-size          0
num-of-cylinders     0
engine-type          0
curb-weight          0
height              0
width               0
length              0
wheel-base          0
engine-location      0

```



```

drive-wheels      0
body-style        0
num-of-doors      0
aspiration        0
fuel-type         0
make              0
normalized-losses 0
symboling         0
dtype: int64

```

#Linear Regression annalyis

```
[23]: df.columns
```

```

[23]: Index(['price', 'highway-mpg', 'city-mpg', 'peak-rpm', 'horsepower',
            'compression-ratio', 'stroke', 'bore', 'fuel-system', 'engine-size',
            'num-of-cylinders', 'engine-type', 'curb-weight', 'height', 'width',
            'length', 'wheel-base', 'engine-location', 'drive-wheels', 'body-style',
            'num-of-doors', 'aspiration', 'fuel-type', 'make', 'normalized-losses',
            'symboling'],
            dtype='object')

```

```
[24]: df.head(20)
```

```

[24]:
   price  highway-mpg  city-mpg  peak-rpm  horsepower  \
0  13495.000000      27      21    5000.0      111.0
1  16500.000000      27      21    5000.0      111.0
2  16500.000000      26      19    5000.0      154.0
3  13950.000000      30      24    5500.0      102.0
4  17450.000000      22      18    5500.0      115.0
5  15250.000000      25      19    5500.0      110.0
6  17710.000000      25      19    5500.0      110.0
7  18920.000000      25      19    5500.0      110.0
8  23875.000000      20      17    5500.0      140.0
9  13207.129353      22      16    5500.0      160.0
10 16430.000000      29      23    5800.0      101.0
11 16925.000000      29      23    5800.0      101.0
12 20970.000000      28      21    4250.0      121.0
13 21105.000000      28      21    4250.0      121.0
14 24565.000000      25      20    4250.0      121.0
15 30760.000000      22      16    5400.0      182.0
16 41315.000000      22      16    5400.0      182.0
17 36880.000000      20      15    5400.0      182.0
18  5151.000000      53      47    5100.0       48.0
19  6295.000000      43      38    5400.0       70.0

   compression-ratio  stroke  bore  fuel-system  engine-size  ...  wheel-base  \

```

0	9.0	2.68	3.47	mpfi	130	...	88.6
1	9.0	2.68	3.47	mpfi	130	...	88.6
2	9.0	3.47	2.68	mpfi	152	...	94.5
3	10.0	3.40	3.19	mpfi	109	...	99.8
4	8.0	3.40	3.19	mpfi	136	...	99.4
5	8.5	3.40	3.19	mpfi	136	...	99.8
6	8.5	3.40	3.19	mpfi	136	...	105.8
7	8.5	3.40	3.19	mpfi	136	...	105.8
8	8.3	3.40	3.13	mpfi	131	...	105.8
9	7.0	3.40	3.13	mpfi	131	...	99.5
10	8.8	2.80	3.50	mpfi	108	...	101.2
11	8.8	2.80	3.50	mpfi	108	...	101.2
12	9.0	3.19	3.31	mpfi	164	...	101.2
13	9.0	3.19	3.31	mpfi	164	...	101.2
14	9.0	3.19	3.31	mpfi	164	...	103.5
15	8.0	3.39	3.62	mpfi	209	...	103.5
16	8.0	3.39	3.62	mpfi	209	...	103.5
17	8.0	3.39	3.62	mpfi	209	...	110.0
18	9.5	3.03	2.91	2bbl	61	...	88.4
19	9.6	3.11	3.03	2bbl	90	...	94.5

	engine-location	drive-wheels	body-style	num-of-doors	aspiration	\
0	front	rwd	convertible	2.0	std	
1	front	rwd	convertible	2.0	std	
2	front	rwd	hatchback	2.0	std	
3	front	fwd	sedan	4.0	std	
4	front	4wd	sedan	4.0	std	
5	front	fwd	sedan	2.0	std	
6	front	fwd	sedan	4.0	std	
7	front	fwd	wagon	4.0	std	
8	front	fwd	sedan	4.0	turbo	
9	front	4wd	hatchback	2.0	turbo	
10	front	rwd	sedan	2.0	std	
11	front	rwd	sedan	4.0	std	
12	front	rwd	sedan	2.0	std	
13	front	rwd	sedan	4.0	std	
14	front	rwd	sedan	4.0	std	
15	front	rwd	sedan	4.0	std	
16	front	rwd	sedan	2.0	std	
17	front	rwd	sedan	4.0	std	
18	front	fwd	hatchback	2.0	std	
19	front	fwd	hatchback	2.0	std	

	fuel-type	make	normalized-losses	symboling
0	gas	alfa-romero	122.0	3
1	gas	alfa-romero	122.0	3
2	gas	alfa-romero	122.0	1

3	gas	audi	164.0	2
4	gas	audi	164.0	2
5	gas	audi	122.0	2
6	gas	audi	158.0	1
7	gas	audi	122.0	1
8	gas	audi	158.0	1
9	gas	audi	122.0	0
10	gas	bmw	192.0	2
11	gas	bmw	192.0	0
12	gas	bmw	188.0	0
13	gas	bmw	188.0	0
14	gas	bmw	122.0	1
15	gas	bmw	122.0	0
16	gas	bmw	122.0	0
17	gas	bmw	122.0	0
18	gas	chevrolet	121.0	2
19	gas	chevrolet	98.0	1

[20 rows x 26 columns]

```
[25]: df.describe()
```

```
[25]:
```

	price	highway-mpg	city-mpg	peak-rpm	horsepower \
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	13207.129353	30.751220	25.219512	5125.369458	104.256158
std	7868.768212	6.886443	6.542142	476.979093	39.519211
min	5118.000000	16.000000	13.000000	4150.000000	48.000000
25%	7788.000000	25.000000	19.000000	4800.000000	70.000000
50%	10595.000000	30.000000	24.000000	5200.000000	95.000000
75%	16500.000000	34.000000	30.000000	5500.000000	116.000000
max	45400.000000	54.000000	49.000000	6600.000000	288.000000

	compression-ratio	stroke	bore	engine-size \
count	205.000000	205.000000	205.000000	205.000000
mean	10.142537	3.255423	3.329751	126.907317
std	3.972040	0.313597	0.270844	41.642693
min	7.000000	2.070000	2.540000	61.000000
25%	8.600000	3.110000	3.150000	97.000000
50%	9.000000	3.290000	3.310000	120.000000
75%	9.400000	3.410000	3.580000	141.000000
max	23.000000	4.170000	3.940000	326.000000

	num-of-cylinders	curb-weight	height	width	length \
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	4.380488	2555.565854	53.724878	65.907805	174.049268
std	1.080854	520.680204	2.443522	2.145204	12.337289
min	2.000000	1488.000000	47.800000	60.300000	141.100000

25%	4.000000	2145.000000	52.000000	64.100000	166.300000
50%	4.000000	2414.000000	54.100000	65.500000	173.200000
75%	4.000000	2935.000000	55.500000	66.900000	183.100000
max	12.000000	4066.000000	59.800000	72.300000	208.100000

	wheel-base	num-of-doors	normalized-losses	symboling
count	205.000000	205.000000	205.000000	205.000000
mean	98.756585	3.123153	122.000000	0.834146
std	6.021776	0.989952	31.681008	1.245307
min	86.600000	2.000000	65.000000	-2.000000
25%	94.500000	2.000000	101.000000	0.000000
50%	97.000000	4.000000	122.000000	1.000000
75%	102.400000	4.000000	137.000000	2.000000
max	120.900000	4.000000	256.000000	3.000000

```
[26]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 205 non-null   float64
1   highway-mpg          205 non-null   int64
2   city-mpg              205 non-null   int64
3   peak-rpm              205 non-null   float64
4   horsepower            205 non-null   float64
5   compression-ratio     205 non-null   float64
6   stroke                205 non-null   float64
7   bore                  205 non-null   float64
8   fuel-system           205 non-null   object
9   engine-size           205 non-null   int64
10  num-of-cylinders       205 non-null   int64
11  engine-type            205 non-null   object
12  curb-weight            205 non-null   int64
13  height                 205 non-null   float64
14  width                  205 non-null   float64
15  length                 205 non-null   float64
16  wheel-base             205 non-null   float64
17  engine-location        205 non-null   object
18  drive-wheels           205 non-null   object
19  body-style             205 non-null   object
20  num-of-doors           205 non-null   float64
21  aspiration             205 non-null   object
22  fuel-type              205 non-null   object
23  make                   205 non-null   object
24  normalized-losses      205 non-null   float64
```

```

25  symboling          205 non-null    int64
dtypes: float64(12), int64(6), object(8)
memory usage: 41.8+ KB

```

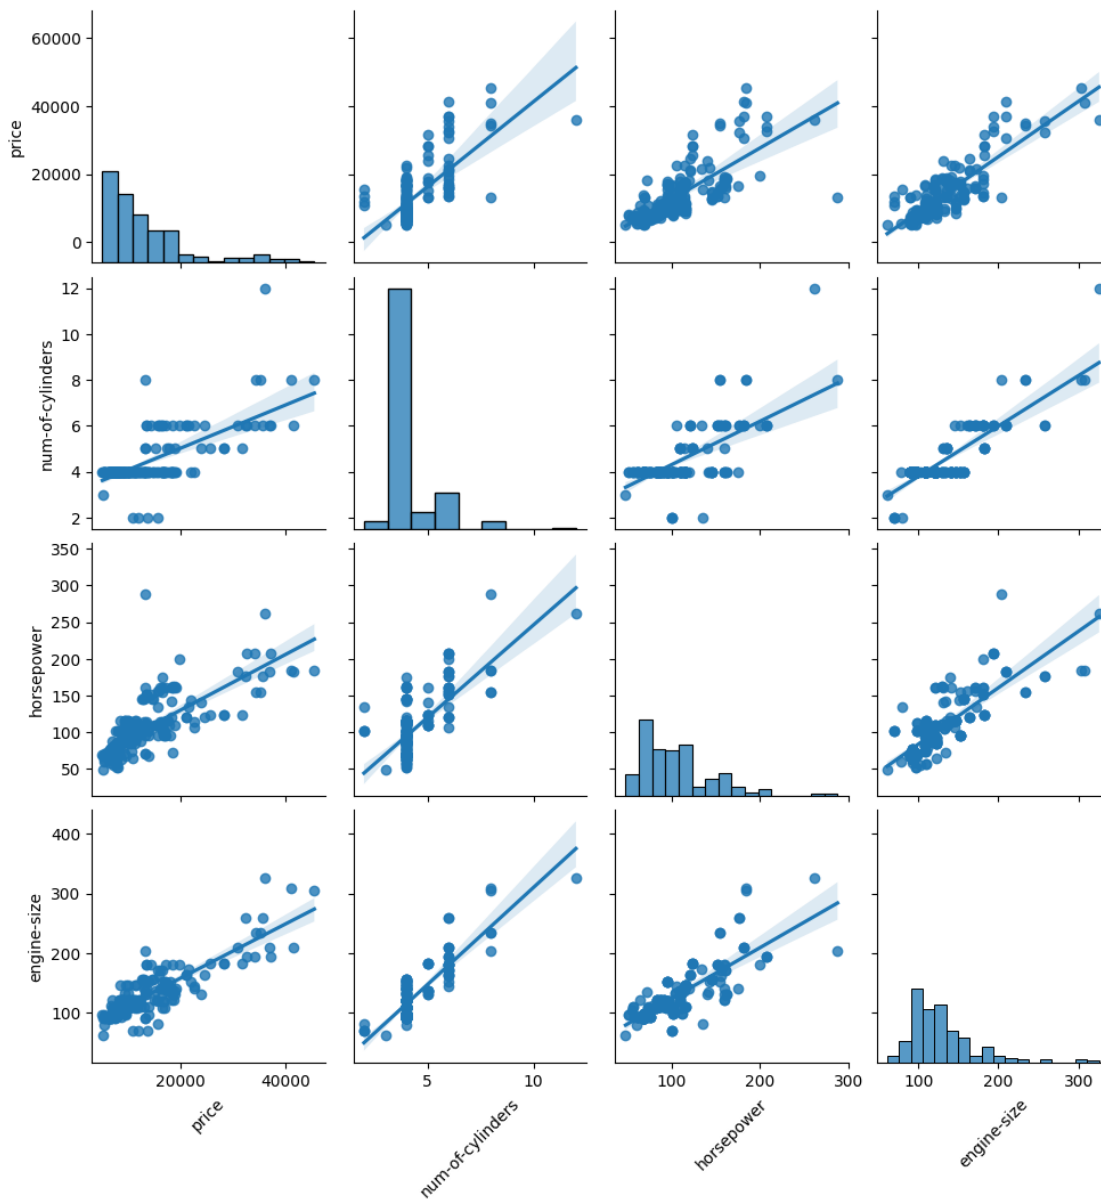
```

[29]: pairplot = sns.pairplot(df[['price', 'num-of-cylinders', 'horsepower',
    ↪ 'engine-size']], kind='reg')

for ax in pairplot.axes.flatten():
    ax.set_xlabel(ax.get_xlabel(), rotation = 45)


plt.show()

```

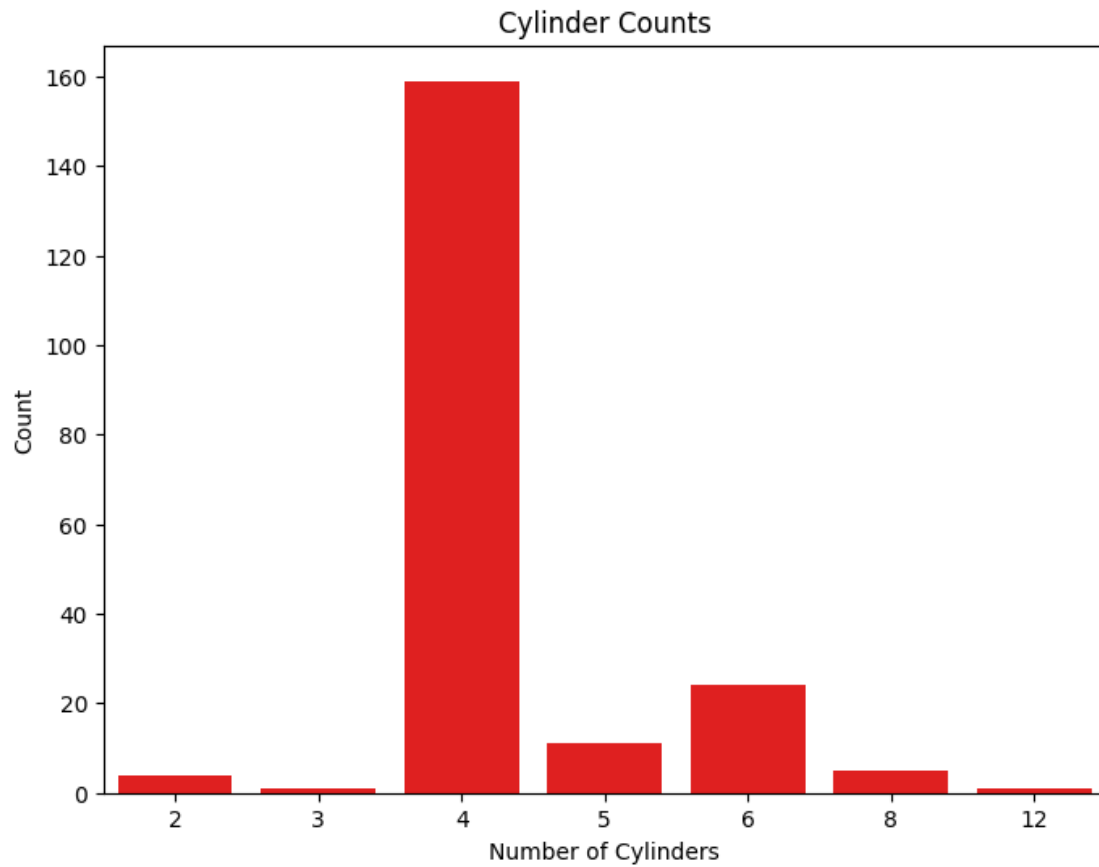


```
[30]: import matplotlib.pyplot as plt

# Plotting the histogram for automobile prices
plt.figure(figsize=(8, 6))
sns.histplot(df['price'], bins=20, kde=True, color='red')
plt.title('Automobile Price')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()
```



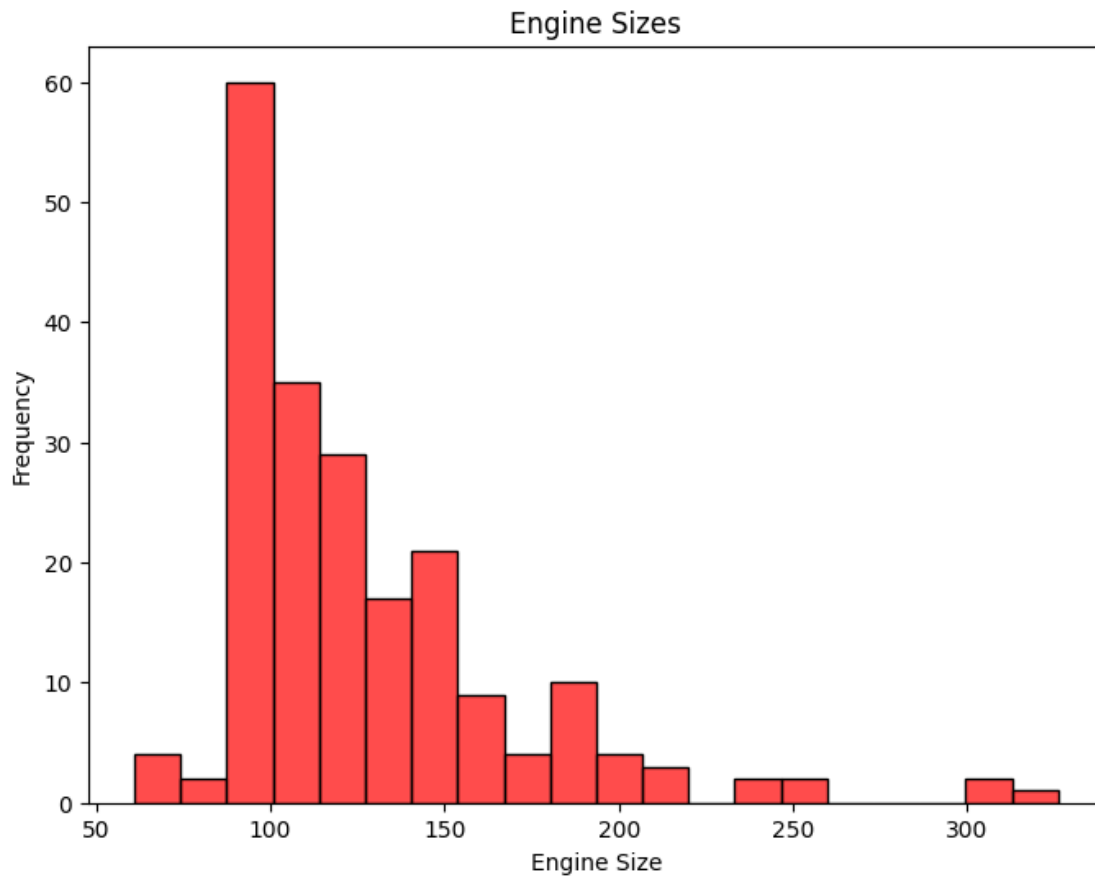
```
[34]: plt.figure(figsize=(8, 6))
sns.countplot(x='num-of-cylinders', color='red', data=df)
plt.title('Cylinder Counts')
plt.xlabel('Number of Cylinders')
plt.ylabel('Count')
plt.show()
```



```
[35]: plt.figure(figsize=(8, 6))
sns.histplot(df['engine-size'], bins=20, color='red', alpha=0.7)


plt.title('Engine Sizes')
plt.xlabel('Engine Size')
plt.ylabel('Frequency')

plt.show()
```




```
[36]: # Create scatterplot for Price vs. Horsepower
plt.figure(figsize=(8, 6))
sns.scatterplot(x='horsepower', y='price', data=df, color='red')
plt.title('Price vs. Horsepower')
plt.xlabel('Horsepower')
plt.ylabel('Price')
plt.show()
```






indicate that cars with higher horsepower tend to be more expensive

```
[37]: # Create scatterplot for Price vs. Engine Size
plt.figure(figsize=(8, 6))
sns.scatterplot(x='engine-size', y='price', data=df, color='red')
plt.title('Price vs. Engine Size')
plt.xlabel('Engine Size')
plt.ylabel('Price')
plt.show()
```



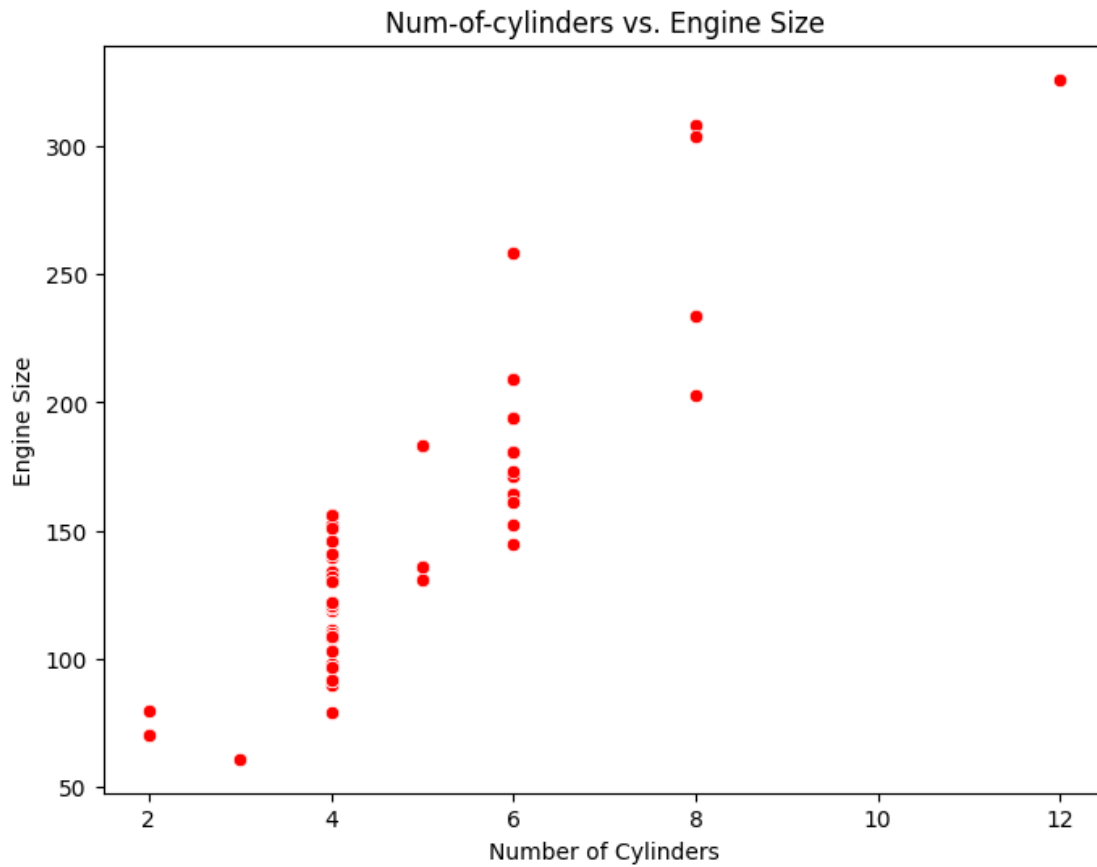
larger engines can be more expensive.

```
[38]: # Create scatterplot for Num-of-cylinders vs. Horsepower
plt.figure(figsize=(8, 6))
sns.scatterplot(x='num-of-cylinders', y='horsepower', data=df, color='red')
plt.title('Num-of-cylinders vs. Horsepower')
plt.xlabel('Number of Cylinders')
plt.ylabel('Horsepower')
plt.show()
```




automobile with more cylinders have higher horsepower.

```
[39]: # Create scatterplot for Num-of-cylinders vs. Engine Size
plt.figure(figsize=(8, 6))
sns.scatterplot(x='num-of-cylinders', y='engine-size', data=df, color='red')
plt.title('Num-of-cylinders vs. Engine Size')
plt.xlabel('Number of Cylinders')
plt.ylabel('Engine Size')
plt.show()
```



manufacturers tend to increase engine size by adding more cylinders

```
[40]: # Create scatterplot for Horsepower vs. Engine Size
plt.figure(figsize=(8, 6))
sns.scatterplot(x='horsepower', y='engine-size', data=df, color='red')
plt.title('Horsepower vs. Engine Size')
plt.xlabel('Horsepower')
plt.ylabel('Engine Size')
plt.show()
```



helps determine that larger engine sizes tend to produce higher horsepower outputs.

```
[44]: print("X =", X.shape, "\nY =", y.shape)
```

```
X = (205, 25)
```

```
Y = (205, 1)
```

```
[76]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

```
[77]: X_train.shape
```


```
[77]: (164, 3)
```

```
[78]: X_test.shape
```

```
[78]: (41, 3)
```

```
[80]: model = LinearRegression()
```

```
[67]: correlation_matrix = df[['price', 'num-of-cylinders', 'horsepower', 'engine-size']].corr()
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix')
plt.show()
```



**##Price vs. num-of-cylinders (0.687770):** There is a moderately positive correlation between the price and the number of cylinders in the engine.

More cylinders often mean larger and more powerful engine, which can result in a higher price for the vehicle.

**##Price vs. horsepower (0.757917):** There is a strong positive correlation between the price and the horsepower of the car.

Higher horsepower is often will be better performance, which can justify a higher price tag for the car.

**##Price vs. engine-size (0.861752):** There is a very strong positive correlation between

the price and the engine size.

Larger engines typically have more power and can offer better performance, leading to higher prices.

**##Num-of-cylinders vs. horsepower (0.691208): There is a moderately positive correlation between the number of cylinders and the horsepower of the car.**

More cylinders generally mean more power generated by the engine, resulting in higher horsepower.

**##Num-of-cylinders vs. engine-size (0.846031): There is a strong positive correlation between the number of cylinders and the engine size.**

More cylinders mean a larger total displacement of the engine, which leads to a larger engine size.

**##Horsepower vs. engine-size (0.810713): There is a strong positive correlation between the horsepower and the engine size.**

A larger engine size allows for more fuel, resulting in higher power output.

```
[84]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LinearRegression
      from sklearn.metrics import mean_squared_error

      X = df[['num-of-cylinders', 'horsepower', 'engine-size']]
      y = df['price']


      y_pred = model.predict(X_test)
      mse = mean_squared_error(y_test, y_pred)
      print("Mean Squared Error:", mse)
      print("Coefficients:", model.coef_)
      print("Intercept:", model.intercept_)
```

Mean Squared Error: 16171180.212878957

Coefficients: [-1201.27854017 37.88655387 157.55524143]


Intercept: -5478.095778018269

```
[92]: # Scatter plot with regression line for price vs num-of-cylinders
      plt.figure(figsize=(6, 4))
      sns.regplot(x='price', y='num-of-cylinders', data=df)
      plt.title('Regression Plot for price vs num-of-cylinders')
      plt.xlabel('price')
      plt.ylabel('num-of-cylinders')
      plt.tight_layout()
      plt.show()
```




```
[93]: # Scatter plot with regression line for price vs horsepower
plt.figure(figsize=(6, 4))
sns.regplot(x='price', y='horsepower', data=df)
plt.title('Regression Plot for price vs horsepower')
plt.xlabel('price')
plt.ylabel('horsepower')
plt.tight_layout()
plt.show()
```





```
[94]: # Scatter plot with regression line for price vs engine-size
plt.figure(figsize=(6, 4))
sns.regplot(x='price', y='engine-size', data=df)
plt.title('Regression Plot for price vs engine-size')
plt.xlabel('price')
plt.ylabel('engine-size')
plt.tight_layout()
plt.show()
```



According to the graphs that the larger and more powerful the automobile tend to be have higher prices.