

# Classification using Logistic Regression

April 28, 2024

## 1 Hands-on Activity 11.2 Classification using Logistic Regression

Course: CPE 311	Program: BSCpE
<b>Course Title:</b> Computational Thinking with Python	<b>Date Performed:</b> April 27 , 2024
<b>Section:</b> BSCPE22S3	<b>Date Submitted:</b> April 28, 2024
<b>Student Name:</b> John Louie V. Adornado	<b>Instructor's Name:</b> Engr. Roman Richard

```
[1]: pip install ucimlrepo
```

```
Collecting ucimlrepo
  Downloading ucimlrepo-0.0.6-py3-none-any.whl (8.0 kB)
Installing collected packages: ucimlrepo
Successfully installed ucimlrepo-0.0.6
```

```
[2]: pip install hvplot
```

```
Collecting hvplot
  Downloading hvplot-0.9.2-py2.py3-none-any.whl (1.8 MB)
    1.8/1.8 MB
18.9 MB/s eta 0:00:00
Requirement already satisfied: bokeh>=1.0.0 in
/usr/local/lib/python3.10/dist-packages (from hvplot) (3.3.4)
Requirement already satisfied: colorcet>=2 in /usr/local/lib/python3.10/dist-
packages (from hvplot) (3.1.0)
Requirement already satisfied: holoviews>=1.11.0 in
/usr/local/lib/python3.10/dist-packages (from hvplot) (1.17.1)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages
(from hvplot) (2.0.3)
Requirement already satisfied: numpy>=1.15 in /usr/local/lib/python3.10/dist-
packages (from hvplot) (1.25.2)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-
packages (from hvplot) (24.0)
Requirement already satisfied: panel>=0.11.0 in /usr/local/lib/python3.10/dist-
packages (from hvplot) (1.3.8)
Requirement already satisfied: param<3.0,>=1.12.0 in
```

/usr/local/lib/python3.10/dist-packages (from hvplot) (2.1.0)  
 Requirement already satisfied: Jinja2>=2.9 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (3.1.3)  
 Requirement already satisfied: contourpy>=1 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (1.2.1)  
 Requirement already satisfied: pillow>=7.1.0 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (9.4.0)  
 Requirement already satisfied: PyYAML>=3.10 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (6.0.1)  
 Requirement already satisfied: tornado>=5.1 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (6.3.3)  
 Requirement already satisfied: xyzservices>=2021.09.1 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (2024.4.0)  
 Requirement already satisfied: pyviz-comms>=0.7.4 in /usr/local/lib/python3.10/dist-packages (from holoviews>=1.11.0->hvplot) (3.0.2)  
 Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->hvplot) (2.8.2)  
 Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->hvplot) (2023.4)  
 Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->hvplot) (2024.1)  
 Requirement already satisfied: markdown in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (3.6)  
 Requirement already satisfied: markdown-it-py in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (3.0.0)  
 Requirement already satisfied: linkify-it-py in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (2.0.3)  
 Requirement already satisfied: mdit-py-plugins in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (0.4.0)  
 Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (2.31.0)  
 Requirement already satisfied: tqdm>=4.48.0 in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (4.66.2)  
 Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (6.1.0)  
 Requirement already satisfied: typing-extensions in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (4.11.0)  
 Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from Jinja2>=2.9->bokeh>=1.0.0->hvplot) (2.1.5)  
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas->hvplot) (1.16.0)  
 Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from bleach->panel>=0.11.0->hvplot) (0.5.1)  
 Requirement already satisfied: uc-micro-py in /usr/local/lib/python3.10/dist-packages (from linkify-it-py->panel>=0.11.0->hvplot) (1.0.3)  
 Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-it-py->panel>=0.11.0->hvplot) (0.1.2)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (3.3.2)  
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (3.7)  
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (2.0.7)  
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (2024.2.2)  
Installing collected packages: hvplot  
Successfully installed hvplot-0.9.2

```
[3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import hvplot.pandas
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.linear_model import LinearRegression
%matplotlib inline
```

```
[4]: from ucimlrepo import fetch_ucirepo

# fetch dataset
cervical_cancer_risk_factors = fetch_ucirepo(id=383)

# data (as pandas dataframes)
X = cervical_cancer_risk_factors.data.features
y = cervical_cancer_risk_factors.data.targets

# metadata
print(cervical_cancer_risk_factors.metadata)

# variable information
print(cervical_cancer_risk_factors.variables)
```

```
{'uci_id': 383, 'name': 'Cervical Cancer (Risk Factors)', 'repository_url':
'https://archive.ics.uci.edu/dataset/383/cervical+cancer+risk+factors',
'data_url': 'https://archive.ics.uci.edu/static/public/383/data.csv',
'abstract': 'This dataset focuses on the prediction of indicators/diagnosis of
cervical cancer. The features cover demographic information, habits, and
historic medical records.', 'area': 'Health and Medicine', 'tasks':
['Classification'], 'characteristics': ['Multivariate'], 'num_instances': 858,
'num_features': 36, 'feature_types': ['Integer', 'Real'], 'demographics':
['Age', 'Other'], 'target_col': None, 'index_col': None, 'has_missing_values':
```

```
'yes', 'missing_values_symbol': 'NaN', 'year_of_dataset_creation': 2017,
'last_updated': 'Sun Mar 10 2024', 'dataset_doi': '10.24432/C5Z310', 'creators':
['Kelwin Fernandes', 'Jaime Cardoso', 'Jessica Fernandes'], 'intro_paper':
{'title': 'Transfer Learning with Partial Observability Applied to Cervical
Cancer Screening', 'authors': 'Kelwin Fernandes, Jaime S. Cardoso, Jessica C.
Fernandes', 'published_in': 'Iberian Conference on Pattern Recognition and Image
Analysis', 'year': 2017, 'url': 'https://www.semanticscholar.org/paper/Transfer-
Learning-with-Partial-Observability-to-Fernandes-
Cardoso/1c02438ba4dfa775399ba414508e9cd335b69012', 'doi': None},
'additional_info': {'summary': "The dataset was collected at 'Hospital
Universitario de Caracas' in Caracas, Venezuela. The dataset comprises
demographic information, habits, and historic medical records of 858 patients.
Several patients decided not to answer some of the questions because of privacy
concerns (missing values).", 'purpose': None, 'funded_by': None,
'instances_represent': None, 'recommended_data_splits': None, 'sensitive_data':
None, 'preprocessing_description': None, 'variable_info': '(int) Age\r\n(int)
Number of sexual partners\r\n(int) First sexual intercourse (age)\r\n(int) Num
of pregnancies\r\n(bool) Smokes\r\n(bool) Smokes (years)\r\n(bool) Smokes
(packs/year)\r\n(bool) Hormonal Contraceptives\r\n(int) Hormonal Contraceptives
(years)\r\n(bool) IUD\r\n(int) IUD (years)\r\n(bool) STDs\r\n(int) STDs
(number)\r\n(bool) STDs:condylomatosis\r\n(bool) STDs:cervical
condylomatosis\r\n(bool) STDs:vaginal condylomatosis\r\n(bool) STDs:vulvo-
perineal condylomatosis\r\n(bool) STDs:syphilis\r\n(bool) STDs:pelvic
inflammatory disease\r\n(bool) STDs:genital herpes\r\n(bool) STDs:molluscum
contagiosum\r\n(bool) STDs:AIDS\r\n(bool) STDs:HIV\r\n(bool) STDs:Hepatitis
B\r\n(bool) STDs:HPV\r\n(int) STDs: Number of diagnosis\r\n(int) STDs: Time
since first diagnosis\r\n(int) STDs: Time since last diagnosis\r\n(bool)
Dx:Cancer\r\n(bool) Dx:CIN\r\n(bool) Dx:HPV\r\n(bool) Dx\r\n(bool) Hinselmann:
target variable\r\n(bool) Schiller: target variable\r\n(bool) Cytology: target
variable\r\n(bool) Biopsy: target variable', 'citation': None}}
```

	name	role	type	demographic \
0	Age	Feature	Integer	Age
1	Number of sexual partners	Feature	Continuous	Other
2	First sexual intercourse	Feature	Continuous	None
3	Num of pregnancies	Feature	Continuous	None
4	Smokes	Feature	Continuous	None
5	Smokes (years)	Feature	Continuous	None
6	Smokes (packs/year)	Feature	Continuous	None
7	Hormonal Contraceptives	Feature	Continuous	None
8	Hormonal Contraceptives (years)	Feature	Continuous	None
9	IUD	Feature	Continuous	None
10	IUD (years)	Feature	Continuous	None
11	STDs	Feature	Continuous	None
12	STDs (number)	Feature	Continuous	None
13	STDs:condylomatosis	Feature	Continuous	None
14	STDs:cervical condylomatosis	Feature	Continuous	None
15	STDs:vaginal condylomatosis	Feature	Continuous	None
16	STDs:vulvo-perineal condylomatosis	Feature	Continuous	None

17	STDs:syphilis	Feature	Continuous	None
18	STDs:pelvic inflammatory disease	Feature	Continuous	None
19	STDs:genital herpes	Feature	Continuous	None
20	STDs:molluscum contagiosum	Feature	Continuous	None
21	STDs:AIDS	Feature	Continuous	None
22	STDs:HIV	Feature	Continuous	None
23	STDs:Hepatitis B	Feature	Continuous	None
24	STDs:HPV	Feature	Continuous	None
25	STDs: Number of diagnosis	Feature	Integer	None
26	STDs: Time since first diagnosis	Feature	Continuous	None
27	STDs: Time since last diagnosis	Feature	Continuous	None
28	Dx:Cancer	Feature	Integer	None
29	Dx:CIN	Feature	Integer	None
30	Dx:HPV	Feature	Integer	None
31	Dx	Feature	Integer	None
32	Hinselmann	Feature	Integer	None
33	Schiller	Feature	Integer	None
34	Citology	Feature	Integer	None
35	Biopsy	Feature	Integer	None

	description	units	missing_values
0	None	None	no
1	None	None	yes
2	None	None	yes
3	None	None	yes
4	None	None	yes
5	None	None	yes
6	None	None	yes
7	None	None	yes
8	None	None	yes
9	None	None	yes
10	None	None	yes
11	None	None	yes
12	None	None	yes
13	None	None	yes
14	None	None	yes
15	None	None	yes
16	None	None	yes
17	None	None	yes
18	None	None	yes
19	None	None	yes
20	None	None	yes
21	None	None	yes
22	None	None	yes
23	None	None	yes
24	None	None	yes
25	None	None	no
26	None	None	yes

27	None	None	yes
28	None	None	no
29	None	None	no
30	None	None	no
31	None	None	no
32	None	None	no
33	None	None	no
34	None	None	no
35	None	None	no

```
[5]: df = pd.concat([X, y], axis = 1)
df
```

```
[5]:      Age  Number of sexual partners  First sexual intercourse \
0      18                        4.0                15.0
1      15                        1.0                14.0
2      34                        1.0                 NaN
3      52                        5.0                16.0
4      46                        3.0                21.0
..    ...
853    34                        3.0                18.0
854    32                        2.0                19.0
855    25                        2.0                17.0
856    33                        2.0                24.0
857    29                        2.0                20.0
```

	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	\
0	1.0	0.0	0.0	0.0	
1	1.0	0.0	0.0	0.0	
2	1.0	0.0	0.0	0.0	
3	4.0	1.0	37.0	37.0	
4	4.0	0.0	0.0	0.0	
..	...	...	...	...	
853	0.0	0.0	0.0	0.0	
854	1.0	0.0	0.0	0.0	
855	0.0	0.0	0.0	0.0	
856	2.0	0.0	0.0	0.0	
857	1.0	0.0	0.0	0.0	

	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	...	\
0	0.0	0.00	0.0	...	
1	0.0	0.00	0.0	...	
2	0.0	0.00	0.0	...	
3	1.0	3.00	0.0	...	
4	1.0	15.00	0.0	...	
..	...	...	...	...	
853	0.0	0.00	0.0	...	

854	1.0	8.00	0.0	...
855	1.0	0.08	0.0	...
856	1.0	0.08	0.0	...
857	1.0	0.50	0.0	...

	STDs: Time since first diagnosis	STDs: Time since last diagnosis	\
0	NaN	NaN	
1	NaN	NaN	
2	NaN	NaN	
3	NaN	NaN	
4	NaN	NaN	
..	...	...	
853	NaN	NaN	
854	NaN	NaN	
855	NaN	NaN	
856	NaN	NaN	
857	NaN	NaN	

	Dx:Cancer	Dx:CIN	Dx:HPV	Dx	Hinselmann	Schiller	Citology	Biopsy
0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0
3	1	0	1	0	0	0	0	0
4	0	0	0	0	0	0	0	0
..	...	...	...	..	...	...	...	
853	0	0	0	0	0	0	0	0
854	0	0	0	0	0	0	0	0
855	0	0	0	0	0	0	1	0
856	0	0	0	0	0	0	0	0
857	0	0	0	0	0	0	0	0

[858 rows x 36 columns]

```
[6]: df.shape
```

```
[6]: (858, 36)
```

```
[7]: df.head()
```

```
[7]:   Age  Number of sexual partners  First sexual intercourse \
0   18                        4.0                    15.0
1   15                        1.0                    14.0
2   34                        1.0                     NaN
3   52                        5.0                    16.0
4   46                        3.0                    21.0
```

Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	\
--------------------	--------	----------------	---------------------	---

0	1.0	0.0	0.0	0.0
1	1.0	0.0	0.0	0.0
2	1.0	0.0	0.0	0.0
3	4.0	1.0	37.0	37.0
4	4.0	0.0	0.0	0.0

	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	...	\
0	0.0		0.0	0.0	...
1	0.0		0.0	0.0	...
2	0.0		0.0	0.0	...
3	1.0		3.0	0.0	...
4	1.0		15.0	0.0	...

	STDs: Time since first diagnosis	STDs: Time since last diagnosis	\
0	NaN		NaN
1	NaN		NaN
2	NaN		NaN
3	NaN		NaN
4	NaN		NaN

	Dx:Cancer	Dx:CIN	Dx:HPV	Dx	Hinselmann	Schiller	Citology	Biopsy
0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0
3	1	0	1	0	0	0	0	0
4	0	0	0	0	0	0	0	0

[5 rows x 36 columns]

```
[8]: df.columns
```

```
[8]: Index(['Age', 'Number of sexual partners', 'First sexual intercourse',
        'Num of pregnancies', 'Smokes', 'Smokes (years)', 'Smokes (packs/year)',
        'Hormonal Contraceptives', 'Hormonal Contraceptives (years)', 'IUD',
        'IUD (years)', 'STDs', 'STDs (number)', 'STDs:condylomatosis',
        'STDs:cervical condylomatosis', 'STDs:vaginal condylomatosis',
        'STDs:vulvo-perineal condylomatosis', 'STDs:syphilis',
        'STDs:pelvic inflammatory disease', 'STDs:genital herpes',
        'STDs:molluscum contagiosum', 'STDs:AIDS', 'STDs:HIV',
        'STDs:Hepatitis B', 'STDs:HPV', 'STDs: Number of diagnosis',
        'STDs: Time since first diagnosis', 'STDs: Time since last diagnosis',
        'Dx:Cancer', 'Dx:CIN', 'Dx:HPV', 'Dx', 'Hinselmann', 'Schiller',
        'Citology', 'Biopsy'],
        dtype='object')
```

```
[9]: df.info()
```



```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 858 entries, 0 to 857
Data columns (total 36 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Age                                         858 non-null    int64
1   Number of sexual partners                 832 non-null    float64
2   First sexual intercourse                  851 non-null    float64
3   Num of pregnancies                        802 non-null    float64
4   Smokes                                     845 non-null    float64
5   Smokes (years)                           845 non-null    float64
6   Smokes (packs/year)                      845 non-null    float64
7   Hormonal Contraceptives                  750 non-null    float64
8   Hormonal Contraceptives (years)          750 non-null    float64
9   IUD                                        741 non-null    float64
10  IUD (years)                              741 non-null    float64
11  STDs                                       753 non-null    float64
12  STDs (number)                            753 non-null    float64
13  STDs:condylomatosis                      753 non-null    float64
14  STDs:cervical condylomatosis             753 non-null    float64
15  STDs:vaginal condylomatosis              753 non-null    float64
16  STDs:vulvo-perineal condylomatosis       753 non-null    float64
17  STDs:syphilis                            753 non-null    float64
18  STDs:pelvic inflammatory disease         753 non-null    float64
19  STDs:genital herpes                      753 non-null    float64
20  STDs:molluscum contagiosum               753 non-null    float64
21  STDs:AIDS                                753 non-null    float64
22  STDs:HIV                                 753 non-null    float64
23  STDs:Hepatitis B                         753 non-null    float64
24  STDs:HPV                                 753 non-null    float64
25  STDs: Number of diagnosis                 858 non-null    int64
26  STDs: Time since first diagnosis          71 non-null     float64
27  STDs: Time since last diagnosis           71 non-null     float64
28  Dx:Cancer                                858 non-null    int64
29  Dx:CIN                                    858 non-null    int64
30  Dx:HPV                                    858 non-null    int64
31  Dx                                         858 non-null    int64
32  Hinselmann                              858 non-null    int64
33  Schiller                                  858 non-null    int64
34  Citology                                  858 non-null    int64
35  Biopsy                                    858 non-null    int64
dtypes: float64(26), int64(10)
memory usage: 241.4 KB

```

```
[10]: df.describe()
```

[10]:

	Age	Number of sexual partners	First sexual intercourse	\			
count	858.000000	832.000000	851.000000				
mean	26.820513	2.527644	16.995300				
std	8.497948	1.667760	2.803355				
min	13.000000	1.000000	10.000000				
25%	20.000000	2.000000	15.000000				
50%	25.000000	2.000000	17.000000				
75%	32.000000	3.000000	18.000000				
max	84.000000	28.000000	32.000000				
	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	\		
count	802.000000	845.000000	845.000000	845.000000			
mean	2.275561	0.145562	1.219721	0.453144			
std	1.447414	0.352876	4.089017	2.226610			
min	0.000000	0.000000	0.000000	0.000000			
25%	1.000000	0.000000	0.000000	0.000000			
50%	2.000000	0.000000	0.000000	0.000000			
75%	3.000000	0.000000	0.000000	0.000000			
max	11.000000	1.000000	37.000000	37.000000			
	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	\			
count	750.000000	750.000000	741.000000				
mean	0.641333	2.256419	0.112011				
std	0.479929	3.764254	0.315593				
min	0.000000	0.000000	0.000000				
25%	0.000000	0.000000	0.000000				
50%	1.000000	0.500000	0.000000				
75%	1.000000	3.000000	0.000000				
max	1.000000	30.000000	1.000000				
...	STDs: Time since first diagnosis	STDs: Time since last diagnosis	\				
count	...	71.000000	71.000000				
mean	...	6.140845	5.816901				
std	...	5.895024	5.755271				
min	...	1.000000	1.000000				
25%	...	2.000000	2.000000				
50%	...	4.000000	3.000000				
75%	...	8.000000	7.500000				
max	...	22.000000	22.000000				
	Dx:Cancer	Dx:CIN	Dx:HPV	Dx	Hinselmann	Schiller	\
count	858.000000	858.000000	858.000000	858.000000	858.000000	858.000000	
mean	0.020979	0.010490	0.020979	0.027972	0.040793	0.086247	
std	0.143398	0.101939	0.143398	0.164989	0.197925	0.280892	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	

75%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

	Citology	Biopsy
count	858.000000	858.000000
mean	0.051282	0.064103
std	0.220701	0.245078
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

[8 rows x 36 columns]

```
[11]: categorical = [var for var in df.columns if df[var].dtype=='O']
print('There are {} categorical variables\n'.format(len(categorical)))
print('The categorical variables are :', categorical)
```

There are 0 categorical variables

The categorical variables are : []

```
[19]: missing_values = df.isnull().sum()
print("Missing Values:")
print(missing_values)
```

```
Missing Values:
Age                                0
Number of sexual partners         26
First sexual intercourse           7
Num of pregnancies                56
Smokes                           13
Smokes (years)                   13
Smokes (packs/year)              13
Hormonal Contraceptives          108
Hormonal Contraceptives (years)  108
IUD                               117
IUD (years)                      117
STDs                             105
STDs (number)                    105
STDs:condylomatosis              105
STDs:cervical condylomatosis     105
STDs:vaginal condylomatosis      105
STDs:vulvo-perineal condylomatosis 105
STDs:syphilis                    105
STDs:pelvic inflammatory disease 105
STDs:genital herpes              105
```

STDs:molluscum contagiosum	105
STDs:AIDS	105
STDs:HIV	105
STDs:Hepatitis B	105
STDs:HPV	105
STDs: Number of diagnosis	0
STDs: Time since first diagnosis	787
STDs: Time since last diagnosis	787
Dx:Cancer	0
Dx:CIN	0
Dx:HPV	0
Dx	0
Hinselmann	0
Schiller	0
Citology	0
Biopsy	0
dtype: int64	

```
[21]: df.fillna(df.median(), inplace=True)
```

```
[22]: df.isnull().sum()
```

```
[22]: Age                                0
      Number of sexual partners          0
      First sexual intercourse           0
      Num of pregnancies                 0
      Smokes                             0
      Smokes (years)                     0
      Smokes (packs/year)                0
      Hormonal Contraceptives            0
      Hormonal Contraceptives (years)    0
      IUD                                0
      IUD (years)                        0
      STDs                               0
      STDs (number)                     0
      STDs:condylomatosis                0
      STDs:cervical condylomatosis        0
      STDs:vaginal condylomatosis         0
      STDs:vulvo-perineal condylomatosis  0
      STDs:syphilis                      0
      STDs:pelvic inflammatory disease    0
      STDs:genital herpes                 0
      STDs:molluscum contagiosum          0
      STDs:AIDS                          0
      STDs:HIV                           0
      STDs:Hepatitis B                   0
      STDs:HPV                           0
```

```

STDs: Number of diagnosis      0
STDs: Time since first diagnosis 0
STDs: Time since last diagnosis 0
Dx:Cancer                     0
Dx:CIN                        0
Dx:HPV                        0
Dx                             0
Hinselmann                   0
Schiller                      0
Citology                     0
Biopsy                       0
dtype: int64

```

```
[53]: print('STDs:HIV contains', len(df['STDs:HIV'].unique()), 'labels')
```

```
STDs:HIV contains 2 labels
```

```
[56]: df['STDs:HIV'].unique()
```

```
[56]: array([0., 1.])
```

```
[58]: df['STDs:HIV'].value_counts()
```

```
[58]: STDs:HIV
0.0    840
1.0     18
Name: count, dtype: int64
```

```
[59]: print('Dx:Cancer contains', len(df['Dx:Cancer'].unique()), 'labels')
```

```
Dx:Cancer contains 2 labels
```

```
[60]: df['Dx:Cancer'].unique()
```

```
[60]: array([0, 1])
```

```
[61]: df['Dx:Cancer'].value_counts()
```

```
[61]: Dx:Cancer
0    840
1     18
Name: count, dtype: int64
```

```
[64]: numerical_columns = df.select_dtypes(include=['int64', 'float64']).columns

numerical_summary = df[numerical_columns].describe().round(2)
print(numerical_summary)
```

	Age	Number of sexual partners	First sexual intercourse \
count	858.00	858.00	858.00
mean	26.82	2.51	17.00
std	8.50	1.64	2.79
min	13.00	1.00	10.00
25%	20.00	2.00	15.00
50%	25.00	2.00	17.00
75%	32.00	3.00	18.00
max	84.00	28.00	32.00

	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year) \
count	858.00	858.00	858.00	858.00
mean	2.26	0.14	1.20	0.45
std	1.40	0.35	4.06	2.21
min	0.00	0.00	0.00	0.00
25%	1.00	0.00	0.00	0.00
50%	2.00	0.00	0.00	0.00
75%	3.00	0.00	0.00	0.00
max	11.00	1.00	37.00	37.00

	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD ... \
count	858.00	858.00	858.0 ...
mean	0.69	2.04	0.1 ...
std	0.46	3.57	0.3 ...
min	0.00	0.00	0.0 ...
25%	0.00	0.00	0.0 ...
50%	1.00	0.50	0.0 ...
75%	1.00	2.00	0.0 ...
max	1.00	30.00	1.0 ...

	STDs: Time since first diagnosis	STDs: Time since last diagnosis \
count	858.00	858.00
mean	4.18	3.23
std	1.79	1.82
min	1.00	1.00
25%	4.00	3.00
50%	4.00	3.00
75%	4.00	3.00
max	22.00	22.00

	Dx:Cancer	Dx:CIN	Dx:HPV	Dx	Hinselmann	Schiller	Citology \
count	858.00	858.00	858.00	858.00	858.00	858.00	858.00
mean	0.02	0.01	0.02	0.03	0.04	0.09	0.05
std	0.14	0.10	0.14	0.16	0.20	0.28	0.22
min	0.00	0.00	0.00	0.00	0.00	0.00	0.00
25%	0.00	0.00	0.00	0.00	0.00	0.00	0.00
50%	0.00	0.00	0.00	0.00	0.00	0.00	0.00
75%	0.00	0.00	0.00	0.00	0.00	0.00	0.00

```
max          1.00    1.00    1.00    1.00          1.00    1.00    1.00
```

```
      Biopsy
count  858.00
mean    0.06
std     0.25
min     0.00
25%     0.00
50%     0.00
75%     0.00
max     1.00
```

```
[8 rows x 36 columns]
```

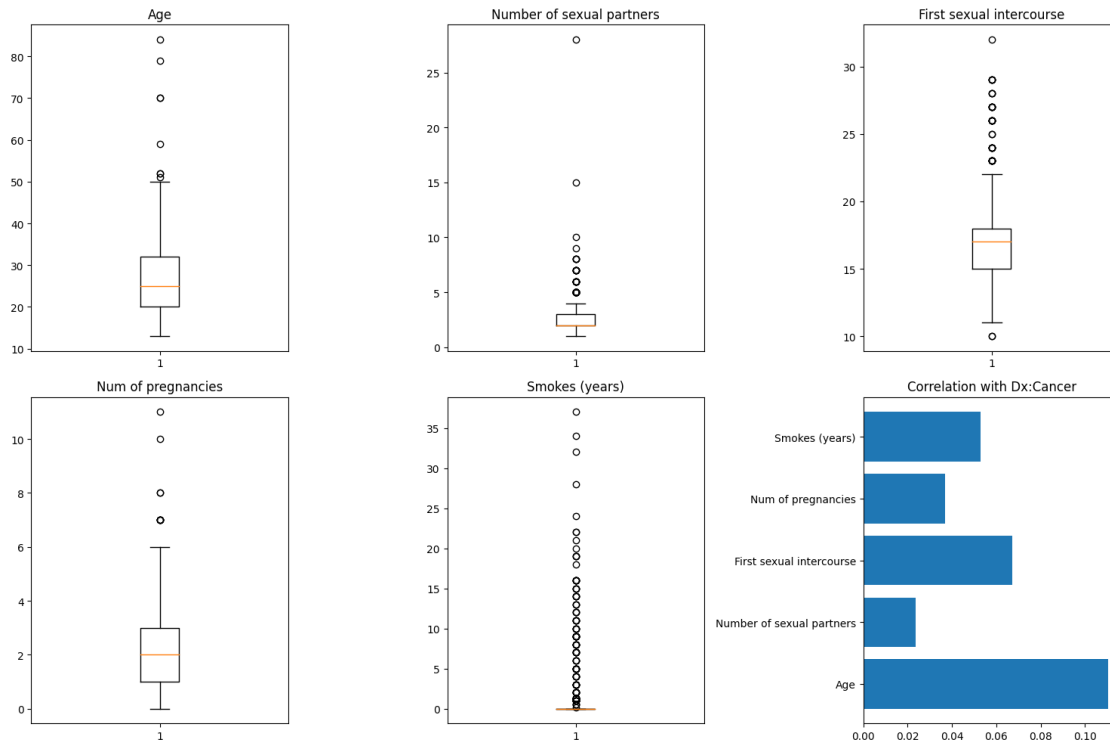
```
[67]: import matplotlib.pyplot as plt

selected_features = ['Age', 'Number of sexual partners', 'First sexual_
↳intercourse',
                    'Num of pregnancies', 'Smokes (years)']

fig, axes = plt.subplots(nrows=2, ncols=3, figsize=(15, 10))


for i, feature in enumerate(selected_features):
    row = i // 3
    col = i % 3
    axes[row, col].boxplot(df[feature].dropna())
    axes[row, col].set_title(feature)

correlation = df[selected_features + ['Dx:Cancer']].corr()['Dx:Cancer']
axes[1, 2].barh(correlation.index[:-1], correlation.values[:-1])
axes[1, 2].set_title('Correlation with Dx:Cancer')
plt.tight_layout()
plt.show()
```

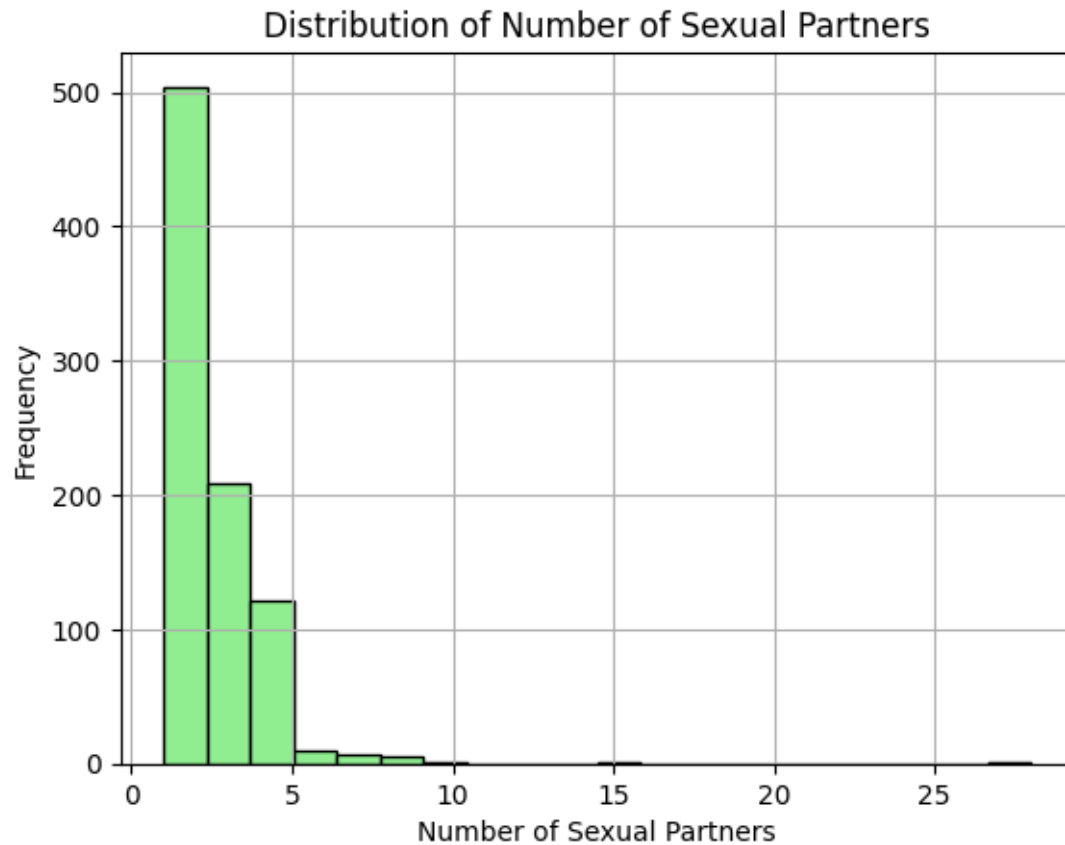


```
[69]: plt.hist(df['Age'].dropna(), bins=20, color='skyblue', edgecolor='black')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.title('Distribution of Age')
plt.grid(True)
plt.show()
```






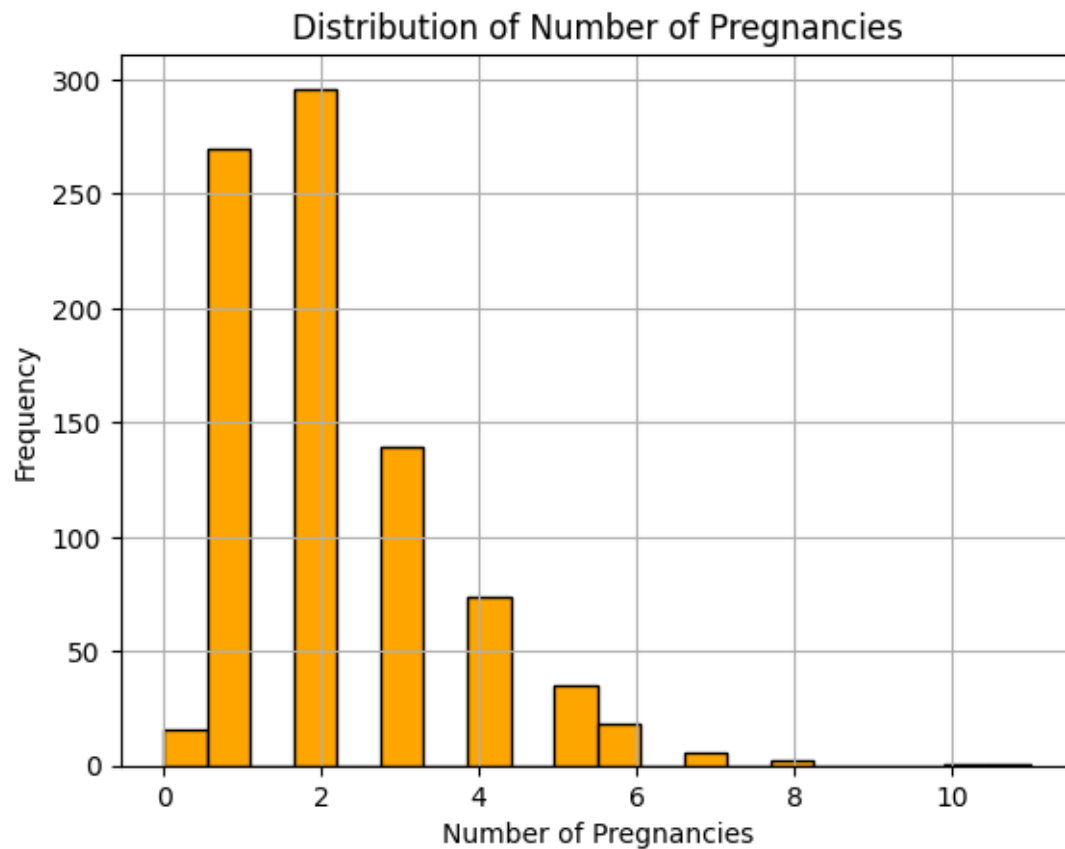
```
[70]: plt.hist(df['Number of sexual partners'].dropna(), bins=20, color='lightgreen',
             edgecolor='black')
plt.xlabel('Number of Sexual Partners')
plt.ylabel('Frequency')
plt.title('Distribution of Number of Sexual Partners')
plt.grid(True)
plt.show()
```



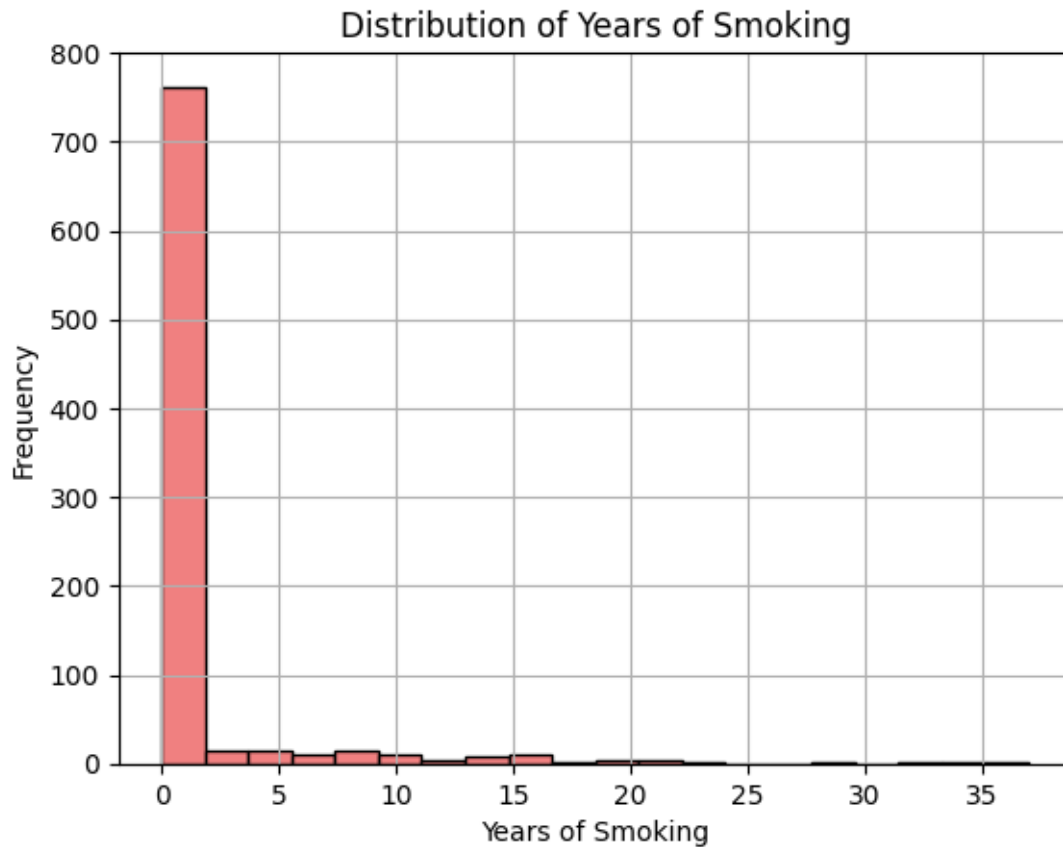
```
[71]: plt.hist(df['First sexual intercourse'].dropna(), bins=20, color='salmon',  
            edgecolor='black')  
plt.xlabel('Age at First Sexual Intercourse')  
plt.ylabel('Frequency')  
plt.title('Distribution of Age at First Sexual Intercourse')  
plt.grid(True)  
plt.show()
```



```
[72]: plt.hist(df['Num of pregnancies'].dropna(), bins=20, color='orange',
             edgecolor='black')
plt.xlabel('Number of Pregnancies')
plt.ylabel('Frequency')
plt.title('Distribution of Number of Pregnancies')
plt.grid(True)
plt.show()
```



```
[73]: plt.hist(df['Smokes (years)'].dropna(), bins=20, color='lightcoral',  
            edgecolor='black')  
plt.xlabel('Years of Smoking')  
plt.ylabel('Frequency')  
plt.title('Distribution of Years of Smoking')  
plt.grid(True)  
plt.show()
```



```
[74]: selected_features = ['Age', 'Number of sexual partners', 'First sexual_
    ↪intercourse',
    ↪'Num of pregnancies', 'Smokes (years)']

outliers = {}
for feature in selected_features:
    q1 = df[feature].quantile(0.25)
    q3 = df[feature].quantile(0.75)
    iqr = q3 - q1
    lower_bound = q1 - 1.5 * iqr
    upper_bound = q3 + 1.5 * iqr
    outliers[feature] = df[(df[feature] < lower_bound) | (df[feature] >_
    ↪upper_bound)][feature]

for feature, outlier_values in outliers.items():
    print(f"Outliers for {feature}:")
    print(outlier_values)
    print()
```

Outliers for Age:

3	52
6	51
652	59
667	79
668	84
671	52
676	70
682	70

Name: Age, dtype: int64

Outliers for Number of sexual partners:

3	5.0
26	5.0
29	6.0
52	5.0
56	5.0

...

796	7.0
804	8.0
812	5.0
837	9.0
844	8.0

Name: Number of sexual partners, Length: 68, dtype: float64

Outliers for First sexual intercourse:

5	23.0
7	26.0
10	26.0
13	25.0
19	27.0
26	23.0
28	24.0
29	26.0
38	24.0
69	23.0
75	27.0
80	27.0
89	32.0
93	29.0
113	23.0
128	29.0
138	28.0
147	10.0
162	23.0
211	26.0
215	24.0
242	23.0

312	27.0
467	10.0
479	24.0
487	26.0
492	26.0
500	29.0
560	24.0
563	28.0
569	25.0
585	28.0
607	29.0
614	23.0
676	27.0
681	29.0
705	23.0
777	26.0
824	27.0
831	23.0
856	24.0

Name: First sexual intercourse, dtype: float64

Outliers for Num of pregnancies:

17	8.0
100	7.0
263	7.0
337	7.0
480	8.0
588	7.0
668	11.0
679	7.0
682	10.0
740	7.0

Name: Num of pregnancies, dtype: float64

Outliers for Smokes (years):

3	37.000000
6	34.000000
9	1.266973
28	3.000000
32	1.266973
...	
834	16.000000
837	11.000000
842	6.000000
844	9.000000
849	11.000000

Name: Smokes (years), Length: 123, dtype: float64

```
[24]: from sklearn.impute import SimpleImputer

imputer = SimpleImputer(strategy='mean')

X_imputed = imputer.fit_transform(X)

X = pd.DataFrame(X_imputed, columns=X.columns)
```

```
[15]: X = df.drop(columns=['Dx:Cancer'])
y = df['Dx:Cancer']
```

```
[38]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

```
[41]: # Initialize a logistic regression model
model = LogisticRegression()

# Fit the model to the training data
model.fit(X_train, y_train)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[41]: LogisticRegression()
```

```
[44]: from sklearn.metrics import classification_report
```

```
# Make predictions on the test set
y_pred = model.predict(X_test_scaled)

# Evaluate the model's performance
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Classification Report
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

```
Accuracy: 0.9825581395348837
```

```
Classification Report:
```



	precision	recall	f1-score	support
0	1.00	0.98	0.99	167
1	0.62	1.00	0.77	5
accuracy			0.98	172
macro avg	0.81	0.99	0.88	172
weighted avg	0.99	0.98	0.98	172

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names

warnings.warn(

```
[76]: from sklearn.model_selection import train_test_split

# Get the minimum and maximum values of features in the training set
X_train_min = X_train.min()
X_train_max = X_train.max()

# Get the minimum and maximum values of features in the testing set
X_test_min = X_test.min()
X_test_max = X_test.max()

# Get the minimum and maximum values of the target variable in the training set
y_train_min = y_train.min()
y_train_max = y_train.max()

# Get the minimum and maximum values of the target variable in the testing set
y_test_min = y_test.min()
y_test_max = y_test.max()

# Print
print("Minimum and maximum values of features in the training set:")
print("X_train_min:", X_train_min)
print("X_train_max:", X_train_max)
print()
print("Minimum and maximum values of features in the testing set:")
print("X_test_min:", X_test_min)
print("X_test_max:", X_test_max)
print()
print("Minimum and maximum values of the target variable in the training set:")
print("y_train_min:", y_train_min)
print("y_train_max:", y_train_max)
print()
print("Minimum and maximum values of the target variable in the testing set:")
print("y_test_min:", y_test_min)
```

```
print("y_test_max:", y_test_max)
```

Minimum and maximum values of features in the training set:

X_train_min: Age	14.0
Number of sexual partners	1.0
First sexual intercourse	10.0
Num of pregnancies	0.0
Smokes	0.0
Smokes (years)	0.0
Smokes (packs/year)	0.0
Hormonal Contraceptives	0.0
Hormonal Contraceptives (years)	0.0
IUD	0.0
IUD (years)	0.0
STDs	0.0
STDs (number)	0.0
STDs:condylomatosis	0.0
STDs:cervical condylomatosis	0.0
STDs:vaginal condylomatosis	0.0
STDs:vulvo-perineal condylomatosis	0.0
STDs:syphilis	0.0
STDs:pelvic inflammatory disease	0.0
STDs:genital herpes	0.0
STDs:molluscum contagiosum	0.0
STDs:AIDS	0.0
STDs:HIV	0.0
STDs:Hepatitis B	0.0
STDs:HPV	0.0
STDs: Number of diagnosis	0.0
STDs: Time since first diagnosis	1.0
STDs: Time since last diagnosis	1.0
Dx:CIN	0.0
Dx:HPV	0.0
Dx	0.0
Hinselmann	0.0
Schiller	0.0
Citology	0.0
Biopsy	0.0
dtype: float64	
X_train_max: Age	79.000000
Number of sexual partners	28.000000
First sexual intercourse	32.000000
Num of pregnancies	10.000000
Smokes	1.000000
Smokes (years)	37.000000
Smokes (packs/year)	37.000000
Hormonal Contraceptives	1.000000

Hormonal Contraceptives (years)	30.000000
IUD	1.000000
IUD (years)	19.000000
STDs	1.000000
STDs (number)	4.000000
STDs:condylomatosis	1.000000
STDs:cervical condylomatosis	0.000000
STDs:vaginal condylomatosis	1.000000
STDs:vulvo-perineal condylomatosis	1.000000
STDs:syphilis	1.000000
STDs:pelvic inflammatory disease	1.000000
STDs:genital herpes	1.000000
STDs:molluscum contagiosum	1.000000
STDs:AIDS	0.000000
STDs:HIV	1.000000
STDs:Hepatitis B	0.001328
STDs:HPV	1.000000
STDs: Number of diagnosis	3.000000
STDs: Time since first diagnosis	22.000000
STDs: Time since last diagnosis	22.000000
Dx:CIN	1.000000
Dx:HPV	1.000000
Dx	1.000000
Hinselmann	1.000000
Schiller	1.000000
Citology	1.000000
Biopsy	1.000000
dtype: float64	

Minimum and maximum values of features in the testing set:

X_test_min: Age	13.0
Number of sexual partners	1.0
First sexual intercourse	11.0
Num of pregnancies	0.0
Smokes	0.0
Smokes (years)	0.0
Smokes (packs/year)	0.0
Hormonal Contraceptives	0.0
Hormonal Contraceptives (years)	0.0
IUD	0.0
IUD (years)	0.0
STDs	0.0
STDs (number)	0.0
STDs:condylomatosis	0.0
STDs:cervical condylomatosis	0.0
STDs:vaginal condylomatosis	0.0
STDs:vulvo-perineal condylomatosis	0.0
STDs:syphilis	0.0

STDs:pelvic inflammatory disease	0.0	
STDs:genital herpes	0.0	
STDs:molluscum contagiosum	0.0	
STDs:AIDS	0.0	
STDs:HIV	0.0	
STDs:Hepatitis B	0.0	
STDs:HPV	0.0	
STDs: Number of diagnosis	0.0	
STDs: Time since first diagnosis	1.0	
STDs: Time since last diagnosis	1.0	
Dx:CIN	0.0	
Dx:HPV	0.0	
Dx	0.0	
Hinselmann	0.0	
Schiller	0.0	
Citology	0.0	
Biopsy	0.0	
dtype: float64		
X_test_max: Age		84.000000
Number of sexual partners	8.000000	
First sexual intercourse	27.000000	
Num of pregnancies	11.000000	
Smokes	1.000000	
Smokes (years)	24.000000	
Smokes (packs/year)	21.000000	
Hormonal Contraceptives	1.000000	
Hormonal Contraceptives (years)	22.000000	
IUD	1.000000	
IUD (years)	12.000000	
STDs	1.000000	
STDs (number)	2.000000	
STDs:condylomatosis	1.000000	
STDs:cervical condylomatosis	0.000000	
STDs:vaginal condylomatosis	0.005312	
STDs:vulvo-perineal condylomatosis	1.000000	
STDs:syphilis	1.000000	
STDs:pelvic inflammatory disease	0.001328	
STDs:genital herpes	0.001328	
STDs:molluscum contagiosum	0.001328	
STDs:AIDS	0.000000	
STDs:HIV	1.000000	
STDs:Hepatitis B	1.000000	
STDs:HPV	1.000000	
STDs: Number of diagnosis	1.000000	
STDs: Time since first diagnosis	21.000000	
STDs: Time since last diagnosis	21.000000	
Dx:CIN	1.000000	
Dx:HPV	1.000000	

```
Dx                                1.000000
Hinselmann                       1.000000
Schiller                         1.000000
Citology                         1.000000
Biopsy                          1.000000
dtype: float64
```

Minimum and maximum values of the target variable in the training set:

```
y_train_min: 0
y_train_max: 1
```

Minimum and maximum values of the target variable in the testing set:

```
y_test_min: 0
y_test_max: 1
```

```
[77]: from sklearn.preprocessing import MinMaxScaler

# Initialize MinMaxScaler
scaler = MinMaxScaler()

# Scale the variables
X_train_scaled = scaler.fit_transform(X_train[['Age', 'Number of sexual_
↪partners', 'First sexual intercourse', 'Num of pregnancies']])
X_test_scaled = scaler.transform(X_test[['Age', 'Number of sexual partners',
↪'First sexual intercourse', 'Num of pregnancies']])
```

```
[81]: from sklearn.preprocessing import StandardScaler

selected_numerical_features = ['Age', 'Number of sexual partners', 'First_
↪sexual intercourse', 'Num of pregnancies']

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train[selected_numerical_features])

X_test_scaled = scaler.transform(X_test[selected_numerical_features])
X_train.describe()
```

```
[81]:
```

	Age	Number of sexual partners	First sexual intercourse \
count	686.000000	686.000000	686.00000
mean	26.666181	2.551218	16.96934
std	8.177105	1.713820	2.78168
min	14.000000	1.000000	10.00000
25%	20.000000	2.000000	15.00000
50%	25.500000	2.000000	17.00000
75%	32.000000	3.000000	18.00000
max	79.000000	28.000000	32.00000

	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	\
count	686.000000	686.000000	686.000000	686.000000	
mean	2.285896	0.137903	1.146960	0.410571	
std	1.379342	0.342143	4.026404	2.149990	
min	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	0.000000	0.000000	0.000000	
50%	2.000000	0.000000	0.000000	0.000000	
75%	3.000000	0.000000	0.000000	0.000000	
max	10.000000	1.000000	37.000000	37.000000	

	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	\
count	686.000000	686.000000	686.000000	
mean	0.644952	2.259937	0.102007	
std	0.446959	3.489038	0.278398	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	
50%	1.000000	1.000000	0.000000	
75%	1.000000	2.445550	0.000000	
max	1.000000	30.000000	1.000000	

	...	STDs: Number of diagnosis	STDs: Time since first diagnosis	\
count	...	686.000000	686.000000	
mean	...	0.090379	6.112902	
std	...	0.311335	1.649027	
min	...	0.000000	1.000000	
25%	...	0.000000	6.140845	
50%	...	0.000000	6.140845	
75%	...	0.000000	6.140845	
max	...	3.000000	22.000000	

	STDs: Time since last diagnosis	Dx:CIN	Dx:HPV	Dx	\
count	686.000000	686.000000	686.000000	686.000000	
mean	5.782819	0.010204	0.018950	0.024781	
std	1.591353	0.100572	0.136449	0.155572	
min	1.000000	0.000000	0.000000	0.000000	
25%	5.816901	0.000000	0.000000	0.000000	
50%	5.816901	0.000000	0.000000	0.000000	
75%	5.816901	0.000000	0.000000	0.000000	
max	22.000000	1.000000	1.000000	1.000000	

	Hinselmann	Schiller	Citology	Biopsy
count	686.000000	686.000000	686.000000	686.000000
mean	0.039359	0.08309	0.056851	0.064140
std	0.194589	0.27622	0.231727	0.245181
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000

50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000

[8 rows x 35 columns]

```
[86]: from sklearn.metrics import accuracy_score
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.9709302325581395

```
[87]: from sklearn.metrics import accuracy_score

y_train_pred = model.predict(X_train_scaled)

train_accuracy = accuracy_score(y_train, y_train_pred)

test_accuracy = accuracy_score(y_test, y_pred)

print("Training set accuracy:", train_accuracy)
print("Testing set accuracy:", test_accuracy)
```

Training set accuracy: 0.9810495626822158

Testing set accuracy: 0.9709302325581395

```
[90]: model_c_001 = LogisticRegression(C=0.01)
model_c_001.fit(X_train_scaled, y_train)
```

[90]: LogisticRegression(C=0.01)

```
[91]: from sklearn.metrics import confusion_matrix

# Get predictions on the testing set
y_pred = model.predict(X_test_scaled)

# Calculate confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)

# Extract values from confusion matrix
tn, fp, fn, tp = conf_matrix.ravel()

# Print confusion matrix and values
print("Confusion Matrix:")
print(conf_matrix)
print()
print("True Negatives:", tn)
```

```
print("False Positives:", fp)
print("False Negatives:", fn)
print("True Positives:", tp)
```

Confusion Matrix:

```
[[167   0]
 [  5   0]]
```

True Negatives: 167

False Positives: 0


False Negatives: 5

True Positives: 0

```
[92]: import seaborn as sns
import matplotlib.pyplot as plt

# Plot confusion matrix heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, cmap="Blues", fmt="d", annot_kws={"size": 16})
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()
```





```
[93]: from sklearn.metrics import classification_report
```

```
# Make predictions on the test set  
y_pred = model.predict(X_test_scaled)  
  
# Evaluate the model's performance  
accuracy = accuracy_score(y_test, y_pred)  
print("Accuracy:", accuracy)  
  
# Classification Report  
print("Classification Report:")  
print(classification_report(y_test, y_pred))
```

Accuracy: 0.9709302325581395

Classification Report:

	precision	recall	f1-score	support
0	0.97	1.00	0.99	167

1	0.00	0.00	0.00	5
accuracy			0.97	172
macro avg	0.49	0.50	0.49	172
weighted avg	0.94	0.97	0.96	172

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
[94]: total_predictions = conf_matrix.sum()
correct_predictions = tp + tn
accuracy = correct_predictions / total_predictions

print("Classification Accuracy:", accuracy)
```

```
Classification Accuracy: 0.9709302325581395
```

```
[95]: misclassified_instances = fp + fn
error = misclassified_instances / total_predictions

print("Classification Error:", error)
```

```
Classification Error: 0.029069767441860465
```

```
[99]: # Compute True Positive Rate (TPR)
tpr = tp / (tp + fn)

print("True Positive Rate (TPR):", tpr)
```

```
True Positive Rate (TPR): 0.0
```

```
[100]: # Compute False Positive Rate (FPR)
fpr = fp / (fp + tn)
```

```
print("False Positive Rate (FPR):", fpr)
```

False Positive Rate (FPR): 0.0

```
[101]: # Compute Specificity (TNR)
specificity = tn / (tn + fp)

print("Specificity (TNR):", specificity)
```

Specificity (TNR): 1.0

```
[103]: probabilities = model.predict_proba(X_test_scaled[:10])

print("Predicted Probabilities for Class 0 and Class 1:")
for i, probs in enumerate(probabilities):
    print(f"Instance {i+1}: Class 0 - {probs[0]}, Class 1 - {probs[1]}")
```

Predicted Probabilities for Class 0 and Class 1:

Instance 1: Class 0 - 0.9938259209521234, Class 1 - 0.006174079047876628  
Instance 2: Class 0 - 0.9872755057306585, Class 1 - 0.012724494269341509  
Instance 3: Class 0 - 0.9804631686006106, Class 1 - 0.01953683139938943  
Instance 4: Class 0 - 0.9832354828923271, Class 1 - 0.01676451710767286  
Instance 5: Class 0 - 0.9917618621438887, Class 1 - 0.008238137856111221  
Instance 6: Class 0 - 0.9913604632347084, Class 1 - 0.00863953676529163  
Instance 7: Class 0 - 0.9930300985664543, Class 1 - 0.006969901433545725  
Instance 8: Class 0 - 0.9892430981437641, Class 1 - 0.010756901856235903  
Instance 9: Class 0 - 0.9848337557236242, Class 1 - 0.015166244276375768  
Instance 10: Class 0 - 0.9674776244569298, Class 1 - 0.03252237554307028

These are the predicted probabilities for the first 10 instances, where Class 0 represents the negative class and Class 1 is high, it suggests that the model believes there is a higher likelihood of cancer.

```
[106]: import pandas as pd

all_probabilities = model.predict_proba(X_test_scaled)

probabilities_df = pd.DataFrame(all_probabilities,
    columns=['Probability_Class_0', 'Probability_Class_1'])

probabilities_df.index = range(1, len(probabilities_df) + 1)

if probabilities_df.loc[1, 'Probability_Class_0'] > probabilities_df.loc[1,
    'Probability_Class_1']:
    title = "Predicted Probabilities - No Cancer (Class 0) and Chance of Cancer_
    (Class 1)"
```

```

else:
    title = "Predicted Probabilities - Chance of No Cancer (Class 0) and Cancer_
↪(Class 1)"

# Display the title

print(title)
print(probabilities_df)

```

Predicted Probabilities - No Cancer (Class 0) and Chance of Cancer (Class 1)

	Probability_Class_0	Probability_Class_1
1	0.993826	0.006174
2	0.987276	0.012724
3	0.980463	0.019537
4	0.983235	0.016765
5	0.991762	0.008238
..	...	...
168	0.989069	0.010931
169	0.987409	0.012591
170	0.987281	0.012719
171	0.984367	0.015633
172	0.988744	0.011256


[172 rows x 2 columns]

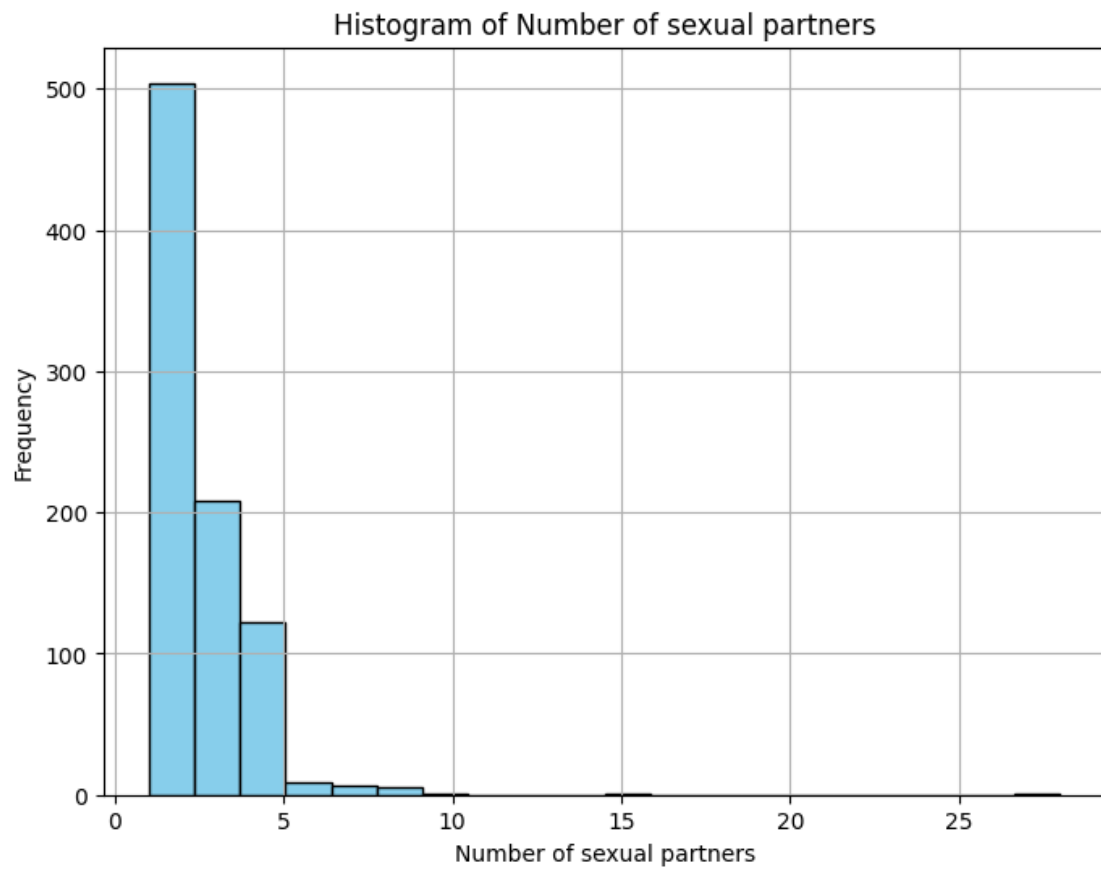
```

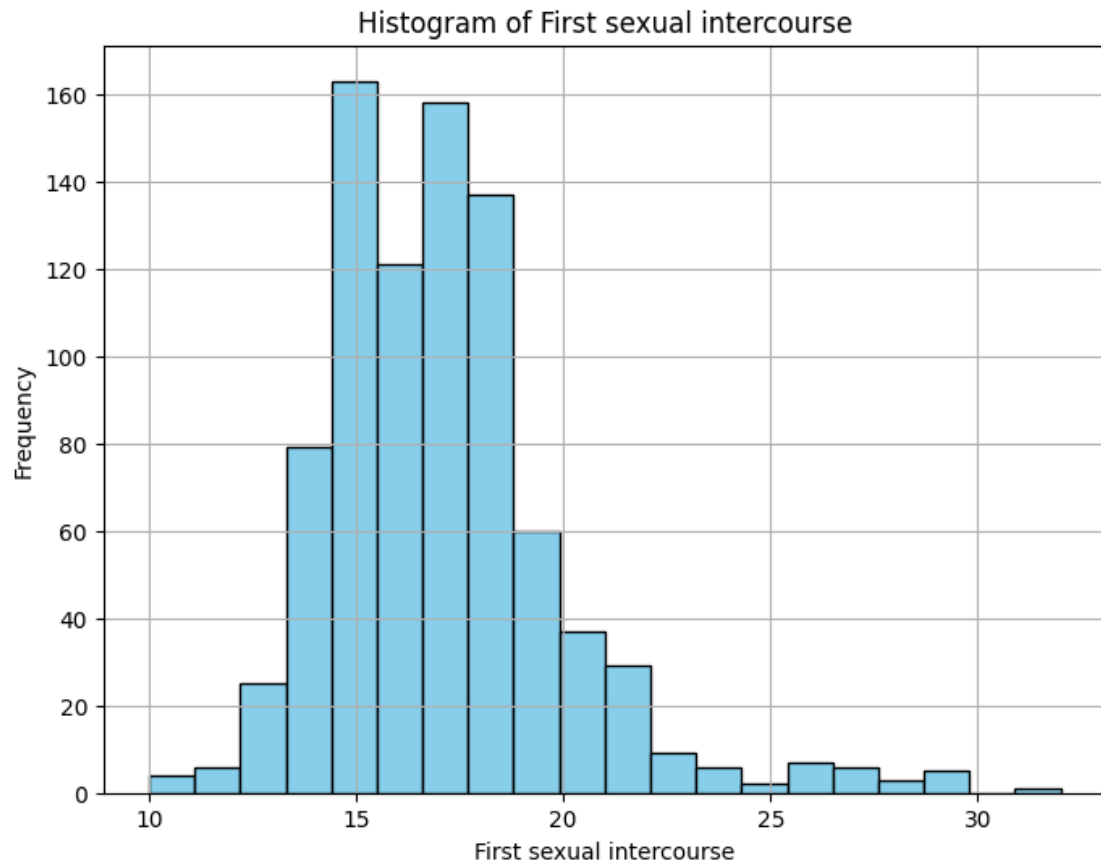
[109]: import matplotlib.pyplot as plt

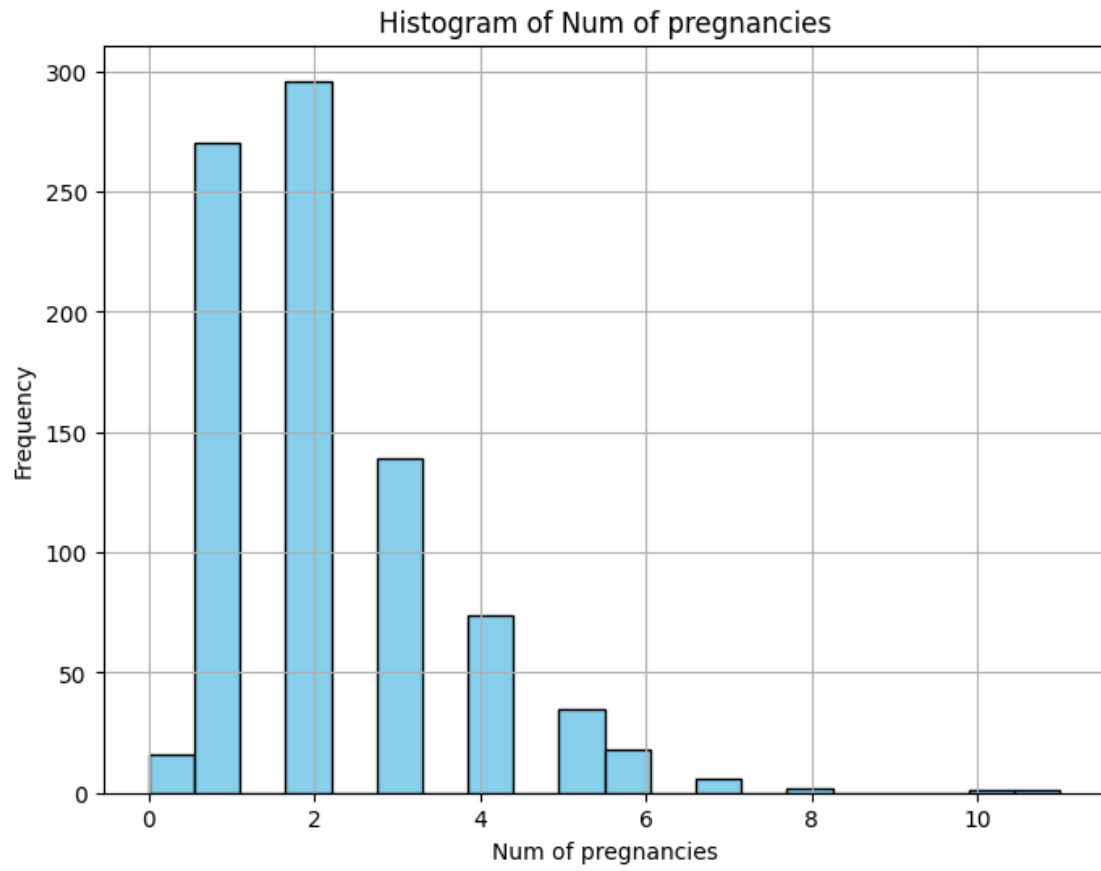
# Plot histograms for each of the five variables
for column in ['Age', 'Number of sexual partners', 'First sexual intercourse',
↪'Num of pregnancies', 'Smokes']:
    plt.figure(figsize=(8, 6))
    plt.hist(df[column].dropna(), bins=20, color='skyblue', edgecolor='black')
    plt.title(f'Histogram of {column}')
    plt.xlabel(column)
    plt.ylabel('Frequency')
    plt.grid(True)
    plt.show()

```

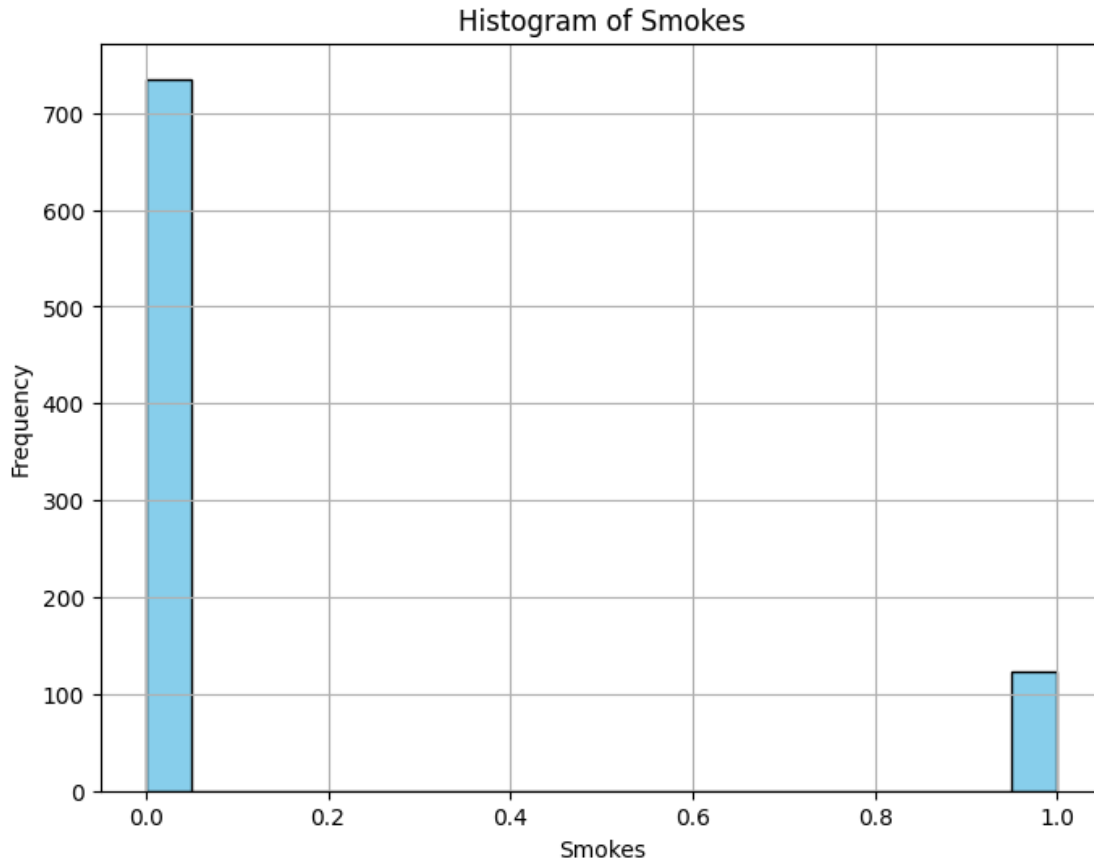












```
[88]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import learning_curve

def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
    ↪n_jobs=None, train_sizes=np.linspace(.1, 1.0, 5)):

    plt.figure()
    plt.title(title)
    if ylim is not None:
        plt.ylim(*ylim)
    plt.xlabel("Training examples")
    plt.ylabel("Score")
    train_sizes, train_scores, test_scores = learning_curve(
        estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)
```

```

plt.grid()


plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                 train_scores_mean + train_scores_std, alpha=0.1,
                 color="r")
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                 test_scores_mean + test_scores_std, alpha=0.1, color="g")
plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
         label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
         label="Cross-validation score")

plt.legend(loc="best")
return plt

# Plot learning curves
title = "Learning Curves (Logistic Regression)"
plot_learning_curve(model, title, X_train_scaled, y_train, cv=5, n_jobs=-1)

plt.show()

```



```
[113]: from sklearn.model_selection import cross_val_score, KFold
from sklearn.linear_model import LogisticRegression
import numpy as np

model = LogisticRegression()

k = 5

kf = KFold(n_splits=k, shuffle=True, random_state=42)

cv_scores = cross_val_score(model, X_train_scaled, y_train, cv=kf,
                             scoring='accuracy')

print("Cross-Validation Scores:", cv_scores)

mean_cv_score = np.mean(cv_scores)
std_cv_score = np.std(cv_scores)
print(f"Mean Cross-Validation Score: {mean_cv_score:.4f}")
print(f"Standard Deviation of Cross-Validation Scores: {std_cv_score:.4f}")
```

Cross-Validation Scores: [0.98550725 0.97810219 0.97080292 0.97810219  
0.99270073]  
Mean Cross-Validation Score: 0.9810  
Standard Deviation of Cross-Validation Scores: 0.0075

```
[116]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression

model = LogisticRegression()

param_grid = {
    'penalty': ['l2'],
    'C': [0.1, 1, 10, 100, 1000]
}

grid_search = GridSearchCV(model, param_grid, cv=5, scoring='accuracy')

grid_search.fit(X_train_scaled, y_train)

best_params = grid_search.best_params_
print("Best Hyperparameters:", best_params)

best_score = grid_search.best_score_
print("Best Cross-Validation Score:", best_score)
```

Best Hyperparameters: {'C': 0.1, 'penalty': 'l2'}

Best Cross-Validation Score: 0.981053633767058

**Comments:**

- Accuracy 98%
- Cross-Validation Scores: 0.0075
- Classification Accuracy: 0.9709302325581395
- Classification Error: 0.029069767441860465
- Specificity (TNR): 1.0

**Summary:**

1. **Model Performance:** The logistic regression model achieved an impressive cross-validation accuracy of approximately 98.11%, indicating its effectiveness in predicting the presence or absence of cervical cancer based on the provided features.
2. **Best Hyperparameters:** The optimal hyperparameters for the logistic regression model were identified as a regularization parameter (C) of 0.1 and an 'l2' penalty.
3. **Significant Features:** Important features contributing to the prediction of cervical cancer include age, number of sexual partners, first sexual intercourse age, and potentially other demographic and clinical factors.
4. **Considerations:** It's important to acknowledge potential limitations such as dataset size, feature selection, and the presence of outliers or missing values, which could impact the model's performance and interpretation of results.

**Overall,** the logistic regression analysis provides valuable insights into the early detection and intervention of cervical cancer.