9sdaccbvu

April 2, 2024

8.2 Querying and Merging

[1]:
```
import pandas as pd
weather = pd.read_csv('nyc_weather_2018.csv')
weather.head()
```

[1]:

	date	datatype	station	attributes	value
0	2018-01-01T00:00:00	PRCP	GHCND:US1CTFR0039	,,N,0800	0.0
1	2018-01-01T00:00:00	PRCP	GHCND:US1NJBG0015	,,N,1050	0.0
2	2018-01-01T00:00:00	SNOW	GHCND:US1NJBG0015	,,N,1050	0.0
3	2018-01-01T00:00:00	PRCP	GHCND:US1NJBG0017	,,N,0920	0.0
4	2018-01-01T00:00:00	SNOW	GHCND:US1NJBG0017	,,N,0920	0.0

Querying DataFrames

[2]:
```
snow_data = weather.query('datatype == "SNOW" and value > 0')
snow_data.head()
```

[2]:

	date	datatype	station	attributes	value
31	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0015	,,N,1600	229.0
34	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0017	,,N,0830	10.0
38	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0018	,,N,0910	46.0
45	2018-01-05T00:00:00	SNOW	GHCND:US1NJBG0017	,,N,0720	102.0
49	2018-01-05T00:00:00	SNOW	GHCND:US1NJBG0018	,,N,1230	183.0

[9]:
```
snow_data2 = weather.loc[(weather["datatype"] == "SNOW") & (weather["value"] >␣
↪0)]
snow_data2.head()
```

[9]:

	date	datatype	station	attributes	value
31	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0015	,,N,1600	229.0
34	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0017	,,N,0830	10.0
38	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0018	,,N,0910	46.0
45	2018-01-05T00:00:00	SNOW	GHCND:US1NJBG0017	,,N,0720	102.0
49	2018-01-05T00:00:00	SNOW	GHCND:US1NJBG0018	,,N,1230	183.0

This is equivalent to quering the data/weather.db SQLite database for SELECT * FROM weather WHERE datatype == “SNOW” AND value $>$ 0 :

[11]:
```
import sqlite3
with sqlite3.connect('weather.db') as connection:
    snow_data_from_db = pd.read_sql('SELECT * FROM weather WHERE datatype ==␣
 ↪"SNOW" AND value > 0', connection )
snow_data.reset_index().drop(columns='index').equals(snow_data_from_db)
```

[11]:
```
True
```

[12]:
```
weather[(weather.datatype == 'SNOW') & (weather.value > 0)].equals(snow_data)
```

[12]:
```
True
```

Merging DataFrames

[13]:
```
station_info = pd.read_csv('weather_stations.csv')
station_info.head()
```

[13]:

	id	name	latitude	longitude	elevation
0	GHCND:US1CTFR0022	STAMFORD 2.6 SSW, CT US	41.064100	-73.577000	36.6
1	GHCND:US1CTFR0039	STAMFORD 4.2 S, CT US	41.037788	-73.568176	6.4
2	GHCND:US1NJBG0001	BERGENFIELD 0.3 SW, NJ US	40.921298	-74.001983	20.1
3	GHCND:US1NJBG0002	SADDLE BROOK TWP 0.6 E, NJ US	40.902694	-74.083358	16.8
4	GHCND:US1NJBG0003	TENAFLY 1.3 W, NJ US	40.914670	-73.977500	21.6

[14]:
```
weather.head()
```

[14]:

	date	datatype	station	attributes	value
0	2018-01-01T00:00:00	PRCP	GHCND:US1CTFR0039	,,N,0800	0.0
1	2018-01-01T00:00:00	PRCP	GHCND:US1NJBG0015	,,N,1050	0.0
2	2018-01-01T00:00:00	SNOW	GHCND:US1NJBG0015	,,N,1050	0.0
3	2018-01-01T00:00:00	PRCP	GHCND:US1NJBG0017	,,N,0920	0.0
4	2018-01-01T00:00:00	SNOW	GHCND:US1NJBG0017	,,N,0920	0.0

to see how many unique values

[15]:
```
station_info.id.describe()
```

[15]:
```
count                   320
unique                  320
top       GHCND:US1CTFR0022
freq                      1
```

```
Name: id, dtype: object
```

[16]:
```
weather.station.describe()
```

[16]:
```
count                  3650
unique                   14
top       GHCND:US1NJBG0017
freq                    576
Name: station, dtype: object
```

[17]:
```
station_info.shape[0], weather.shape[0]
```

[17]:
```
(320, 3650)
```

[18]:
```
def get_row_count(*dfs):
  return [df.shape[0] for df in dfs]
get_row_count(station_info, weather)
```

[18]:
```
[320, 3650]
```

[19]:
```
def get_info(attr, *dfs):
  return list(map(lambda x: getattr(x, attr), dfs))
get_info('shape', station_info, weather)
```

[19]:
```
[(320, 5), (3650, 5)]
```

[20]:
```
inner_join = weather.merge(station_info, left_on='station', right_on='id')
inner_join.sample(5, random_state=0)
```

[20]:

	date	datatype	station	attributes	value	\
3516	2018-09-26T00:00:00	PRCP	GHCND:US1NJBG0010	,,N,0800	84.8	
569	2018-05-23T00:00:00	PRCP	GHCND:US1NJBG0015	,,N,1120	17.0	
991	2018-02-28T00:00:00	DAPR	GHCND:US1NJBG0017	,,N,0750	8.0	
2095	2018-02-22T00:00:00	SNOW	GHCND:US1NJBG0023	,,N,0800	0.0	
2566	2018-08-22T00:00:00	PRCP	GHCND:US1NJBG0030	,,N,0700	21.1	

	id	name	latitude	longitude	\
3516	GHCND:US1NJBG0010	RIVER VALE TWP 1.5 S, NJ US	40.991450	-74.012348	
569	GHCND:US1NJBG0015	NORTH ARLINGTON 0.7 WNW, NJ US	40.791492	-74.139790	
991	GHCND:US1NJBG0017	GLEN ROCK 0.7 SSE, NJ US	40.951090	-74.118264	
2095	GHCND:US1NJBG0023	OAKLAND 0.9 SSE, NJ US	41.019050	-74.233383	
2566	GHCND:US1NJBG0030	OAKLAND 1.0 ESE, NJ US	41.025324	-74.223632	

	elevation
3516	9.4
569	17.7
991	28.0

```
2095      149.4
2566      109.4
```

We can remove the duplication of information in the station and id columns by renaming one of them before the merge and then simply using on

[21]:
```
weather.merge(station_info.rename(dict(id='station'), axis=1), on='station').
  ↪sample(5, random_state=0)
```

[21]:
```
                    date datatype            station attributes  value  \
3516  2018-09-26T00:00:00     PRCP  GHCND:US1NJBG0010   ,,N,0800   84.8
569   2018-05-23T00:00:00     PRCP  GHCND:US1NJBG0015   ,,N,1120   17.0
991   2018-02-28T00:00:00     DAPR  GHCND:US1NJBG0017   ,,N,0750    8.0
2095  2018-02-22T00:00:00     SNOW  GHCND:US1NJBG0023   ,,N,0800    0.0
2566  2018-08-22T00:00:00     PRCP  GHCND:US1NJBG0030   ,,N,0700   21.1

                               name   latitude  longitude  elevation
3516     RIVER VALE TWP 1.5 S, NJ US  40.991450 -74.012348        9.4
569   NORTH ARLINGTON 0.7 WNW, NJ US  40.791492 -74.139790       17.7
991         GLEN ROCK 0.7 SSE, NJ US  40.951090 -74.118264       28.0
2095          OAKLAND 0.9 SSE, NJ US  41.019050 -74.233383      149.4
2566          OAKLAND 1.0 ESE, NJ US  41.025324 -74.223632      109.4
```

[22]:
```
left_join = station_info.merge(weather, left_on='id', right_on='station',␣
  ↪how='left')
right_join = weather.merge(station_info, left_on='station', right_on='id',␣
  ↪how='right')
right_join.tail()
```

[22]:
```
     date datatype station attributes  value                 id  \
3951  NaN      NaN     NaN        NaN    NaN  GHCND:USW00054787
3952  NaN      NaN     NaN        NaN    NaN  GHCND:USW00094728
3953  NaN      NaN     NaN        NaN    NaN  GHCND:USW00094741
3954  NaN      NaN     NaN        NaN    NaN  GHCND:USW00094745
3955  NaN      NaN     NaN        NaN    NaN  GHCND:USW00094789

                                    name  latitude  longitude  elevation
3951  FARMINGDALE REPUBLIC AIRPORT, NY US  40.73443  -73.41637       22.8
3952          NY CITY CENTRAL PARK, NY US  40.77898  -73.96925       42.7
3953             TETERBORO AIRPORT, NJ US  40.85898  -74.05616        0.8
3954        WESTCHESTER CO AIRPORT, NY US  41.06236  -73.70454      112.9
3955     JFK INTERNATIONAL AIRPORT, NY US  40.63915  -73.76390        2.7
```

[23]:
```
left_join.sort_index(axis=1).sort_values(['date', 'station']).reset_index().
  ↪drop(columns='index').equals(
right_join.sort_index(axis=1).sort_values(['date', 'station']).reset_index().
  ↪drop(columns='index')
```

```
)
```

[23]:
```
True
```

[24]:
```
get_info('shape', inner_join, left_join, right_join)
```

[24]:
```
[(3650, 10), (3956, 10), (3956, 10)]
```

[25]:
```
outer_join = weather.merge(
station_info[station_info.name.str.contains('NY')],
left_on='station', right_on='id', how='outer', indicator=True
)
outer_join.sample(4, random_state=0).append(outer_join[outer_join.station.
 ↪isna()].head(2))
```

```
<ipython-input-25-81b63e73e04e>:5: FutureWarning: The frame.append method is
deprecated and will be removed from pandas in a future version. Use
pandas.concat instead.
  outer_join.sample(4,
random_state=0).append(outer_join[outer_join.station.isna()].head(2))
```

[25]:

	date	datatype	station	attributes	value	\
538	2018-05-02T00:00:00	SNOW	GHCND:US1NJBG0015	,,N,0815	0.0	
526	2018-04-23T00:00:00	SNOW	GHCND:US1NJBG0015	,,N,1015	0.0	
2215	2018-05-20T00:00:00	PRCP	GHCND:US1NJBG0023	,,N,0745	12.4	
2872	2018-03-01T00:00:00	SNOW	GHCND:US1NJBG0003	,,N,0730	0.0	
3650	NaN	NaN	NaN	NaN	NaN	
3651	NaN	NaN	NaN	NaN	NaN	

	id	name	latitude	longitude	\
538	NaN	NaN	NaN	NaN	
526	NaN	NaN	NaN	NaN	
2215	NaN	NaN	NaN	NaN	
2872	NaN	NaN	NaN	NaN	
3650	GHCND:US1NJHD0002	KEARNY 1.7 NW, NJ US	40.772892	-74.140926	
3651	GHCND:US1NJHD0018	KEARNY 1.7 NNW, NJ US	40.774342	-74.137109	

	elevation	_merge
538	NaN	left_only
526	NaN	left_only
2215	NaN	left_only
2872	NaN	left_only
3650	29.0	right_only
3651	25.6	right_only

These joins are equivalent to their SQL counterparts. Below is the inner join. Note that to use equals() you will have to do some manipulation of the dataframes to line them up

[27]:
```
import sqlite3
with sqlite3.connect('weather.db') as connection:
  inner_join_from_db = pd.read_sql('SELECT * FROM weather JOIN stations ON␣
 ↪weather.station == stations.id',connection)
inner_join_from_db.shape == inner_join.shape
```

[27]:
```
True
```

[29]:
```
dirty_data = pd.read_csv(
'dirty_data.csv', index_col='date'
).drop_duplicates().drop(columns='SNWD')
dirty_data.head()
```

[29]:
```
                            station  PRCP   SNOW    TMAX  TMIN  TOBS  WESF  \
date
2018-01-01T00:00:00               ?   0.0    0.0  5505.0 -40.0   NaN   NaN
2018-01-02T00:00:00  GHCND:USC00280907   0.0    0.0    -8.3 -16.1 -12.2   NaN
2018-01-03T00:00:00  GHCND:USC00280907   0.0    0.0    -4.4 -13.9 -13.3   NaN
2018-01-04T00:00:00               ?  20.6  229.0  5505.0 -40.0   NaN  19.3
2018-01-05T00:00:00               ?   0.3    NaN  5505.0 -40.0   NaN   NaN

                    inclement_weather
date
2018-01-01T00:00:00               NaN
2018-01-02T00:00:00             False
2018-01-03T00:00:00             False
2018-01-04T00:00:00              True
2018-01-05T00:00:00               NaN
```

[30]:
```
valid_station = dirty_data.query('station != "?"').copy().drop(columns=['WESF',␣
 ↪'station'])
station_with_wesf = dirty_data.query('station == "?"').copy().
 ↪drop(columns=['station', 'TOBS', 'TMIN', 'TMAX'])
```

[31]:
```
valid_station.merge(
station_with_wesf, left_index=True, right_index=True
).query('WESF > 0').head()
```

[31]:
```
                     PRCP_x  SNOW_x  TMAX  TMIN  TOBS inclement_weather_x  \
date
2018-01-30T00:00:00     0.0     0.0   6.7  -1.7  -0.6               False
2018-03-08T00:00:00    48.8     NaN   1.1  -0.6   1.1               False
2018-03-13T00:00:00     4.1    51.0   5.6  -3.9   0.0                True
2018-03-21T00:00:00     0.0     0.0   2.8  -2.8   0.6               False
2018-04-02T00:00:00     9.1   127.0  12.8  -1.1  -1.1                True

                     PRCP_y  SNOW_y  WESF inclement_weather_y
```

```
date
2018-01-30T00:00:00     1.5    13.0   1.8               True
2018-03-08T00:00:00    28.4     NaN  28.7                NaN
2018-03-13T00:00:00     3.0    13.0   3.0               True
2018-03-21T00:00:00     6.6   114.0   8.6               True
2018-04-02T00:00:00    14.0   152.0  15.2               True
```

[32]:
```
valid_station.merge(
station_with_wesf, left_index=True, right_index=True, suffixes=('', '_?')
).query('WESF > 0').head()
```

[32]:
```
                     PRCP   SNOW  TMAX  TMIN  TOBS inclement_weather  PRCP_?  \
date
2018-01-30T00:00:00   0.0    0.0   6.7  -1.7  -0.6             False     1.5
2018-03-08T00:00:00  48.8    NaN   1.1  -0.6   1.1             False    28.4
2018-03-13T00:00:00   4.1   51.0   5.6  -3.9   0.0              True     3.0
2018-03-21T00:00:00   0.0    0.0   2.8  -2.8   0.6             False     6.6
2018-04-02T00:00:00   9.1  127.0  12.8  -1.1  -1.1              True    14.0

                     SNOW_?  WESF inclement_weather_?
date
2018-01-30T00:00:00    13.0   1.8                True
2018-03-08T00:00:00     NaN  28.7                 NaN
2018-03-13T00:00:00    13.0   3.0                True
2018-03-21T00:00:00   114.0   8.6                True
2018-04-02T00:00:00   152.0  15.2                True
```

[33]:
```
valid_station.join(station_with_wesf, rsuffix='_?').query('WESF > 0').head()
```

[33]:
```
                     PRCP   SNOW  TMAX  TMIN  TOBS inclement_weather  PRCP_?  \
date
2018-01-30T00:00:00   0.0    0.0   6.7  -1.7  -0.6             False     1.5
2018-03-08T00:00:00  48.8    NaN   1.1  -0.6   1.1             False    28.4
2018-03-13T00:00:00   4.1   51.0   5.6  -3.9   0.0              True     3.0
2018-03-21T00:00:00   0.0    0.0   2.8  -2.8   0.6             False     6.6
2018-04-02T00:00:00   9.1  127.0  12.8  -1.1  -1.1              True    14.0

                     SNOW_?  WESF inclement_weather_?
date
2018-01-30T00:00:00    13.0   1.8                True
2018-03-08T00:00:00     NaN  28.7                 NaN
2018-03-13T00:00:00    13.0   3.0                True
2018-03-21T00:00:00   114.0   8.6                True
2018-04-02T00:00:00   152.0  15.2                True
```

[34]:
```
weather.set_index('station', inplace=True)
station_info.set_index('id', inplace=True)
```

[35]:
```
weather.index.intersection(station_info.index)
```

[35]:
```
Index(['GHCND:US1CTFR0039', 'GHCND:US1NJBG0015', 'GHCND:US1NJBG0017',
       'GHCND:US1NJBG0018', 'GHCND:US1NJBG0023', 'GHCND:US1NJBG0030',
       'GHCND:US1NJBG0039', 'GHCND:US1NJBG0003', 'GHCND:US1NJBG0044',
       'GHCND:US1NJES0018', 'GHCND:US1NJBG0010', 'GHCND:US1NJES0019',
       'GHCND:US1NJES0024', 'GHCND:US1NJBG0037'],
      dtype='object')
```

[36]:
```
weather.index.difference(station_info.index)
```

[36]:
```
Index([], dtype='object')
```

[37]:
```
station_info.index.difference(weather.index)
```

[37]:
```
Index(['GHCND:US1CTFR0022', 'GHCND:US1NJBG0001', 'GHCND:US1NJBG0002',
       'GHCND:US1NJBG0005', 'GHCND:US1NJBG0006', 'GHCND:US1NJBG0008',
       'GHCND:US1NJBG0011', 'GHCND:US1NJBG0012', 'GHCND:US1NJBG0013',
       'GHCND:US1NJBG0020',
       …
       'GHCND:USW00014708', 'GHCND:USW00014732', 'GHCND:USW00014734',
       'GHCND:USW00014786', 'GHCND:USW00054743', 'GHCND:USW00054787',
       'GHCND:USW00094728', 'GHCND:USW00094741', 'GHCND:USW00094745',
       'GHCND:USW00094789'],
      dtype='object', length=306)
```

[38]:
```
ny_in_name = station_info[station_info.name.str.contains('NY')]
ny_in_name.index.difference(weather.index).shape[0]\
+ weather.index.difference(ny_in_name.index).shape[0]\
== weather.index.symmetric_difference(ny_in_name.index).shape[0]
```

[38]:
```
True
```

[42]:
```
weather.index.unique().union(station_info.index)
```

[42]:
```
Index(['GHCND:US1CTFR0022', 'GHCND:US1CTFR0039', 'GHCND:US1NJBG0001',
       'GHCND:US1NJBG0002', 'GHCND:US1NJBG0003', 'GHCND:US1NJBG0005',
       'GHCND:US1NJBG0006', 'GHCND:US1NJBG0008', 'GHCND:US1NJBG0010',
       'GHCND:US1NJBG0011',
       …
       'GHCND:USW00014708', 'GHCND:USW00014732', 'GHCND:USW00014734',
       'GHCND:USW00014786', 'GHCND:USW00054743', 'GHCND:USW00054787',
       'GHCND:USW00094728', 'GHCND:USW00094741', 'GHCND:USW00094745',
       'GHCND:USW00094789'],
      dtype='object', length=320)
```

[41]:
```
ny_in_name = station_info[station_info.name.str.contains('NY')]
ny_in_name.index.difference(weather.index).union(weather.index.
  ↪difference(ny_in_name.index)).equals(
weather.index.symmetric_difference(ny_in_name.index)
)
```

[41]:
```
True
```